

# CS614: Linux Kernel Programming

## Logistics, Introduction

Debadatta Mishra, CSE, IIT Kanpur

# Course and instructors

\$whereis cs614

Mon, Wed 2PM - 3.15PM RM-101

<https://www.cse.iitk.ac.in/users/deba/cs614/>

Piazza link: <https://piazza.com/iitk.ac.in/secondsemester2023/cs614>

Canvas: <https://canvas.cse.iitk.ac.in/login>

\$whereis deba

KD 212, [deba@cse.iitk.ac.in](mailto:deba@cse.iitk.ac.in) , Meeting hours: 3PM - 5PM Thursday

\$ ls TAs

Arun KP ([kparun@cse.iitk.ac.in](mailto:kparun@cse.iitk.ac.in)), Rohit Singh ([rohit@cse.iitk.ac.in](mailto:rohit@cse.iitk.ac.in))

# Course policy

## Add/Drop

Course registration and drop:- Ideally before 16th Jan 2023

## Class guidelines

Keep your mobile phones switched off / silent

Ask questions and interact (both in class and in Piazza)

Be on time!

# Evaluation

1. Quizzes (10%)
2. Assignments (40%)
3. Mid-semester (15%)
4. End-semester (35%)

# Evaluation

1. Quizzes (10%)
2. Assignments (40%)
3. Mid-semester (15%)
4. End-semester (35%)

## References

Operating Systems: Three Easy Pieces. Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau.

Understanding the Linux Kernel, Daniel P. Bovet, Marco Cesati.

Linux Kernel Development, 3rd Edition, Robert Love.

Linux Device Drivers, 3rd Edition, By Jonathan Corbet, Greg Kroah-Hartman, Alessandro Rubini.

Linux kernel documentation, Research papers

# Evaluation

1. Quizzes (10%)
2. Assignments (40%)
3. Mid-semester (15%)
4. End-semester (35%)

“Take pride in honest hard work”

“Cheating implies accepting defeat”

“If you are here to learn, never defeat the purpose by cheating”

<https://www.cse.iitk.ac.in/pages/AntiCheatingPolicy.html>

# Homework -1

- Personal laptops with decent backup desirable
- **HW1: Setup a Virtual machine for the course (Due before next lecture)**
  - Create a Linux VM (Ubuntu Linux recommended) (KVM is preferred)
  - Download the Linux kernel version - linux-6.1.4
  - Compile and boot the latest kernel

# OS refresher: True/False

1. Superuser (e.g., root user in UNIX) in a multi-tasking OS can execute all instructions provided by the hardware instruction set architecture (ISA).
2. Every process in a computer system is guaranteed to be in running state at least once during its life time.
3. A critical section consisting of a single instruction may require mutual exclusion.
4. A user process interrupted by a device interrupt is always scheduled immediately after the interrupt is serviced.
5. A page fault can be handled without changing any page table entry.



# OS refresher: Quiz

In a uniprocessor system, in which of the following case(s), a process executing in user mode can cause an entry into the OS?

- A. accessing a general purpose register like RAX
- B. executing a JMP (jump) instruction
- C. decrementing an unsigned integer value stored in a register beyond zero
- D. executing a printf( ) statement
- E. returning from a function

# Revisiting Hello World!

hello.c

```
int main(void)
{
    printf("Hello world");
    while(1);
}
```

- What will be the execution behavior and why?
- Alternate ways to print “hello world” on screen?

# Program execution




You said only CPU can execute!

# Inside program execution




You said only CPU can execute!

## CPU execution (from hardware perspective)



- Loads instruction pointed to by PC
- Decode instruction
- Load operand into registers
- Execute instruction (ALU)
- Store results

# Role of the OS



- OS bridges the *semantic gap* between the notions of application execution and real execution
  - OS loads an executable from disk to memory, allocates/frees memory dynamically
  - OS initializes the CPU state i.e., the PC and other registers
  - OS provides interfaces to access I/O devices
- OS facilitates hardware resource sharing and management (How?)

# Virtual view of resources

- Process
  - Each running process thinks that it owns the CPU

# Virtual view of resources

- Process
  - Each running process thinks that it owns the CPU
- Address space
  - Each process feels like it has a huge address space

# Virtual view of resources

- Process
  - Each running process thinks that it owns the CPU
- Address space
  - Each process feels like it has a huge address space
- File system tree
  - The user feels like operating on the files directly



# Virtual view of resources

- Process
  - Each running process thinks that it owns the CPU
- Address space
  - Each process feels like it has a huge address space
- File system tree
  - The user feels like operating on the files directly
- What are the OS responsibilities in providing the above virtual notions?

# Virtual view of resources

- Process
  - Each running process thinks that it owns the CPU
- Address space
  - Each process feels like it has a huge address space
- File system tree
  - The user feels like operating on the files directly
- What are the OS responsibilities in providing the above virtual notions?
  - The OS performs multiplexing of physical resources efficiently
  - Maintains mapping of virtual view to physical resource

# Virtualization: Efficiency/performance

- Resource virtualization should not add *excessive* overheads
- Efficient when programs use the resources directly, infrequent OS mediation
  - Example: when a process is scheduled on CPU, it should execute without OS intervention
- What is the catch?

# Virtualization: Efficiency/performance

- Resource virtualization should not add *excessive* overheads
- Efficient when programs use the resources directly, infrequent OS mediation
  - Example: when a process is scheduled on CPU, it should execute without OS intervention
- What is the catch?
  - Loss of control e.g., process running an infinite loop on a CPU
  - Isolation issues e.g., process accessing/changing OS data structures

# Virtualization: Efficiency/performance

- Resource virtualization should not add *excessive* overheads
- Efficient when programs use the resources directly, infrequent OS mediation
  - Example: when a process is scheduled on CPU, it should execute without OS intervention
- What is the catch?
  - Loss of control e.g., process running an infinite loop on a CPU
  - Isolation issues e.g., process accessing/changing OS data structures

Conclusion: Some limits to direct access must be enforced.

# Limited direct execution

- Can the OS enforce limits to an executing process by itself?


# A process in execution

I want to take control of the CPU from this process which is executing an infinite loop, but how?

OS



I want to restrict this process accessing memory of other processes, but how?  
Monitoring each memory access is not efficient!



# A process in execution


I want to take control of the CPU from this process which is executing an infinite loop, but how?

Help me!

OS



I want to restrict this process accessing memory of other processes, but how?  
Monitoring each memory access is not efficient!





# Limited direct execution

- Can the OS enforce limits to an executing process by itself?
- No, the OS can not enforce limits by itself and still achieve efficiency
- OS requires support from hardware!


# Limited direct execution

- Can the OS enforce limits to an executing process?
- No, the OS can not enforce limits by itself and still achieve efficiency
- OS requires support from hardware!
- What kind of support is needed from the hardware?

# Hardware support: Privilege levels




Help me!




CPU

Hey OS! You can use my support for different privilege levels. Using this you can restrict some operations by the user process!

# Hardware support: Privilege levels




# Hardware support: Privilege levels




- CPU can execute in two modes: *user-mode* and *kernel-mode*
- Some operations are allowed only from kernel-mode (privileged OPs)
  - If executed from user mode, hardware will notify the OS by raising a fault/trap

# Hardware support: Privilege levels




# Hardware support: Privilege levels



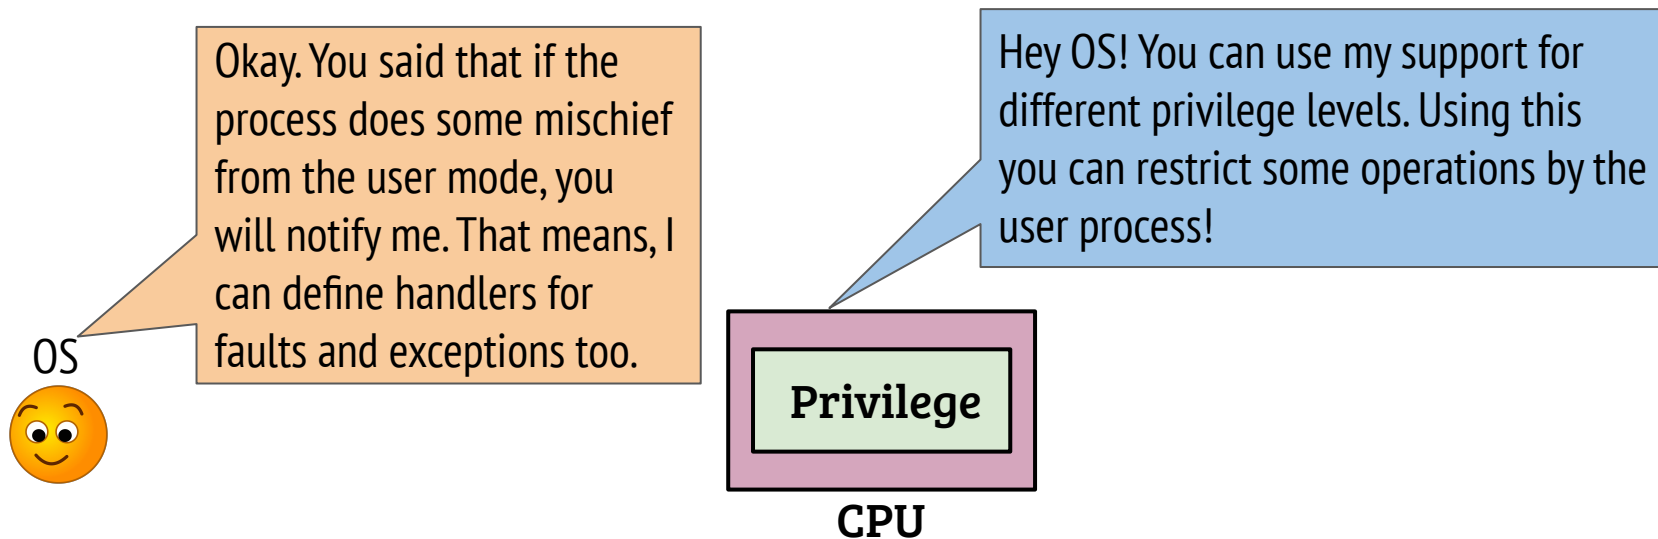
- From user-mode, privilege level of CPU can not be changed directly
- The hardware provides entry instructions from the user-mode which causes a mode switch
- The OS can define the handler for different entry gates

# Hardware support: Privilege levels






# Hardware support: Privilege levels




- The OS can register the handlers for faults and exceptions
- The OS can also register handlers for device interrupts
- *Registration of handlers is privileged!*

# Hardware support: Privilege levels

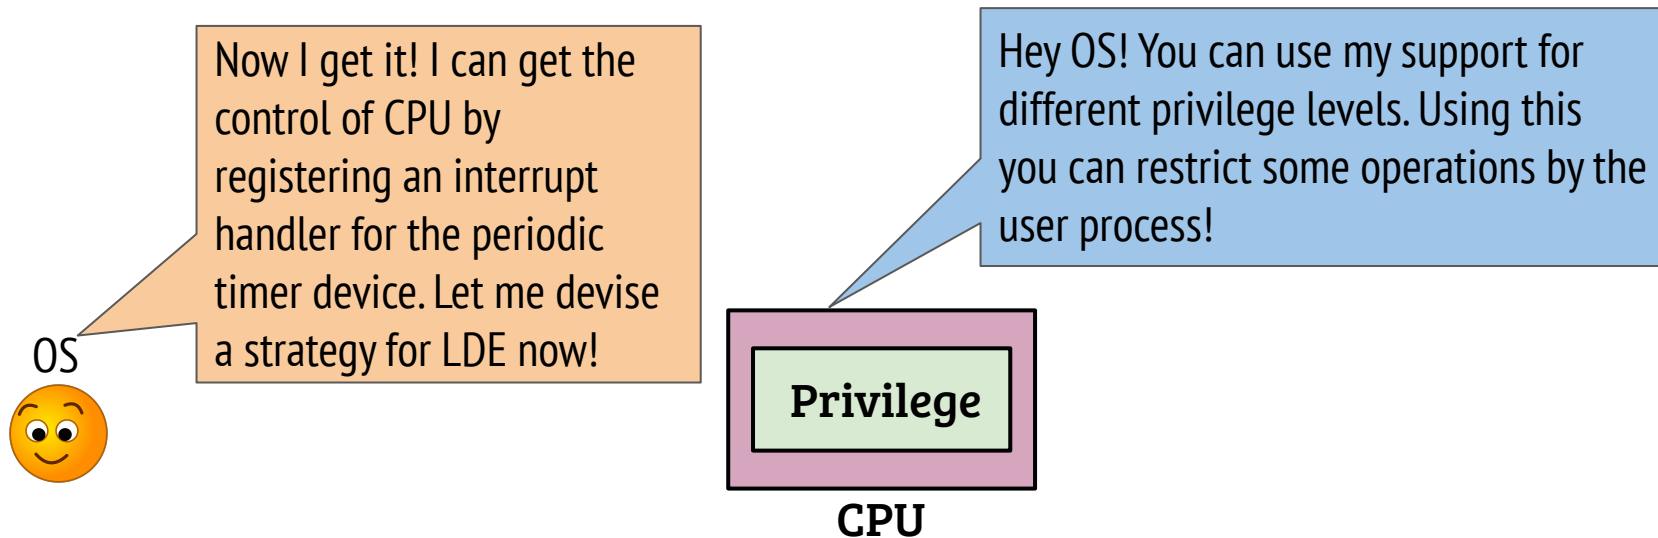


# Hardware support: Privilege levels



- After the boot, the OS needs to configure the handlers for system calls, exceptions/faults and interrupts

# Hardware support: Privilege levels



- After the boot, the OS needs to configure the handlers for system calls, exceptions/faults and interrupts
- The handler code is invoked by the OS when user-mode process invokes a system call or an exception or an external interrupt

# Limited direct execution

- Can the OS enforce limits to an executing process?
- No, the OS can not enforce limits by itself and still achieve efficiency
- OS requires support from hardware!
- What kind of support is needed from the hardware?
- CPU privilege levels: user-mode vs. kernel-mode
- Switching between modes, entry points and handlers

# Evidence based Proof of LDE!

- “Proof” is a term associated with formal world
- “Evidence based proof” is very important for this course
- Proving LDE in the Linux kernel
  - How to prove using a user program executing an infinite loop?

# Evidence based Proof of LDE!

- “Proof” is a term associated with formal world
- “Evidence based proof” is very important for this course
- Proving LDE in the Linux kernel
  - How to prove using a user program executing an infinite loop?
  - CPU usage in user mode should dominate
  - Is this evidence enough?

# Evidence based Proof of LDE!

- “Proof” is a term associated with formal world
- “Evidence based proof” is very important for this course
- Proving LDE in the Linux kernel
  - How to prove using a user program executing an infinite loop?
  - CPU usage in user mode should dominate
  - Is this evidence enough? OS intervention yet to be shown!
  - If a program does division by zero infinitely, we can prove OS intervention. How to go about it?



# Evidence based Proof of LDE!

- “Proof” is a term associated with formal world
- “Evidence based proof” is very important for this course
- Proving LDE in the Linux kernel
  - How to prove using a user program executing an infinite loop?
  - CPU usage in user mode should dominate
  - Is this evidence enough? OS intervention yet to be shown!
  - If a program does division by zero infinitely, we can prove OS intervention. How to go about it?
  - Handle the signal and ignore it or something better (modify division-by-zero handler in OS)  $\Rightarrow$  First hand evidence

# CS614: Linux Kernel Programming


## Privileges and Execution Contexts

Debadatta Mishra, CSE, IIT Kanpur

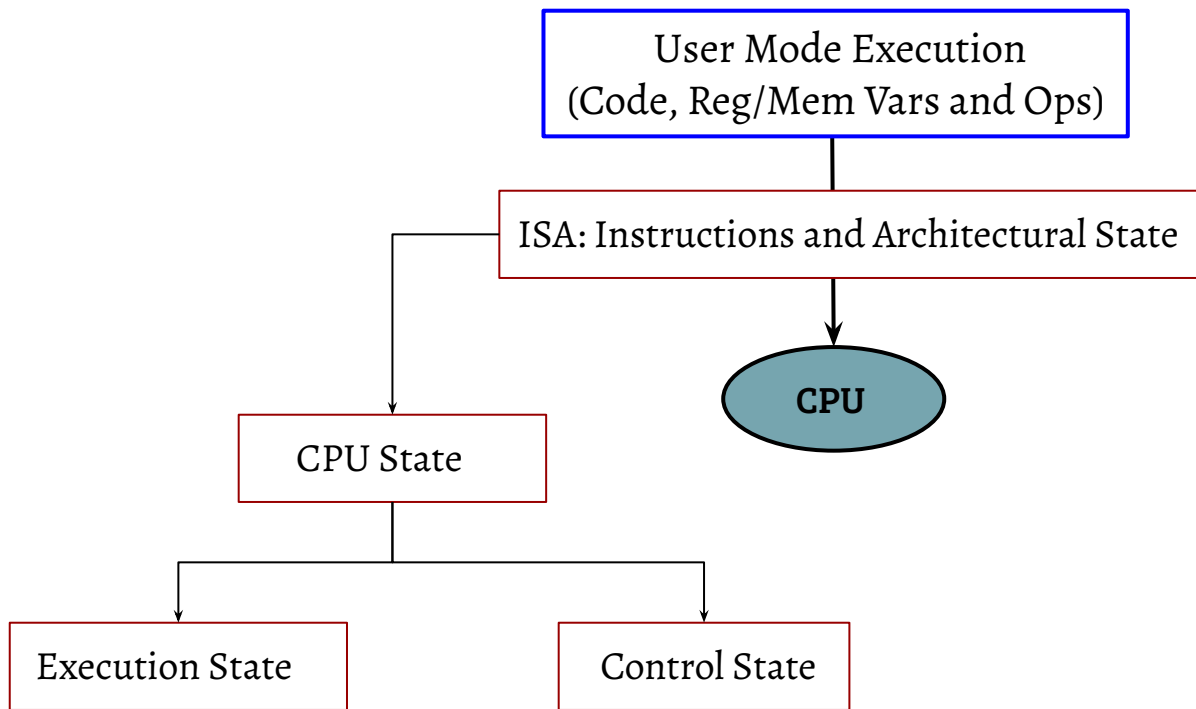
# Recap: Limited direct execution

- Can the OS enforce limits to an executing process?
- No, the OS can not enforce limits by itself and still achieve efficiency
- OS requires support from hardware!
- What kind of support is needed from the hardware?
- CPU privilege levels: user-mode vs. kernel-mode
- Switching between modes, entry points and handlers

# User Mode Execution: ISA and hardware resources



# User Mode Execution: ISA and hardware resources

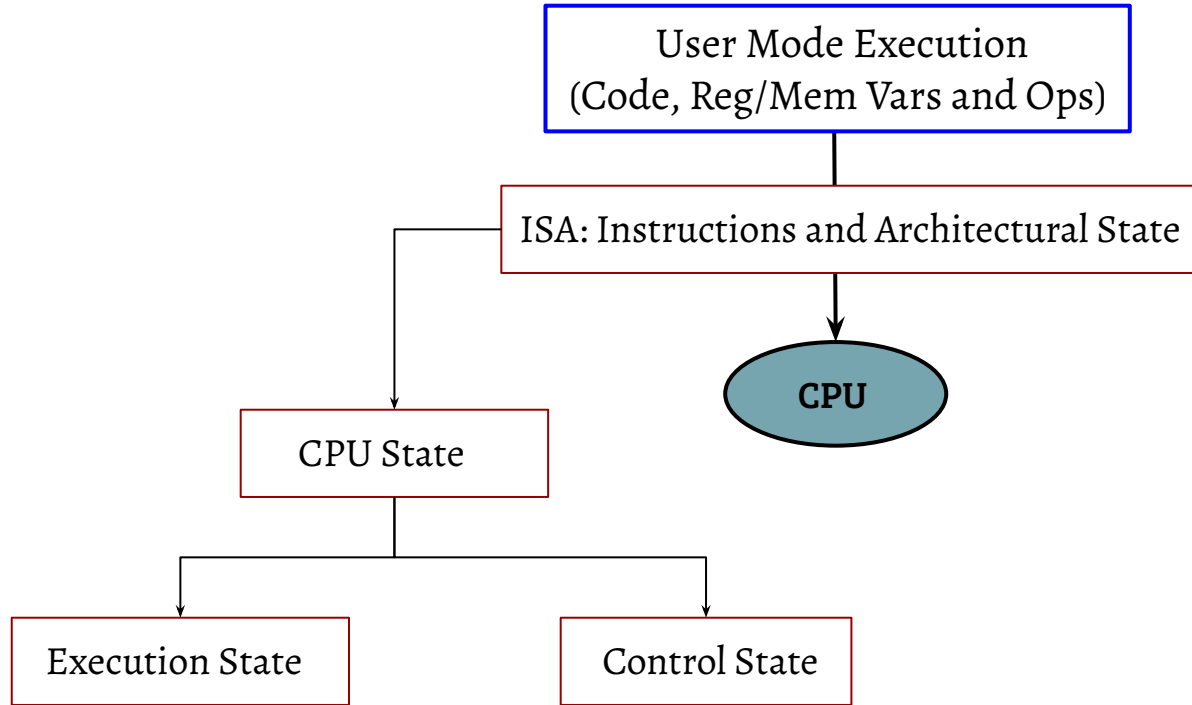


- What is the OS role to ensure correct user mode execution?
- What about memory state? Is the stack a memory state or a register state?

General purpose registers and special registers (IP, SP etc.)

Control registers dictating the CPU behavior (e.g., CRs in X86)

# User Mode Execution: ISA and hardware resources




General purpose registers and special registers (IP, SP etc.)

Control registers dictating the CPU behavior (e.g., CRs in X86)

- What is the OS role to ensure correct user mode execution? OS intervention should not mess-up the state
- What about memory state? Is the stack a memory state or a register state? Stack is maintained in memory, accessed using SP

# X86: rings of protection



- 4 privilege levels: 0 → highest, 3 → lowest
- Some operations are allowed only in privilege level 0
- Most OSes use 0 (for kernel) and 3 (for user)
- Different kinds of privilege enforcement
  - Instruction is privileged
  - Operand is privileged

# Privileged instruction: HLT (on Linux x86\_64)

```
int main( )  
{  
    asm("hlt;");  
}
```

- HLT: Halt the CPU core till next external interrupt
- Executed from user space results in protection fault
- Action: Linux kernel kills the application




# Privileged operation: Read CR3 (Linux x86\_64)

```
#include<stdio.h>
int main(){
    unsigned long cr3_val;
    asm volatile("mov %%cr3, %0;"
                : "=r" (cr3_val)
                :: );
    printf("%lx\n", cr3_val);
}
```

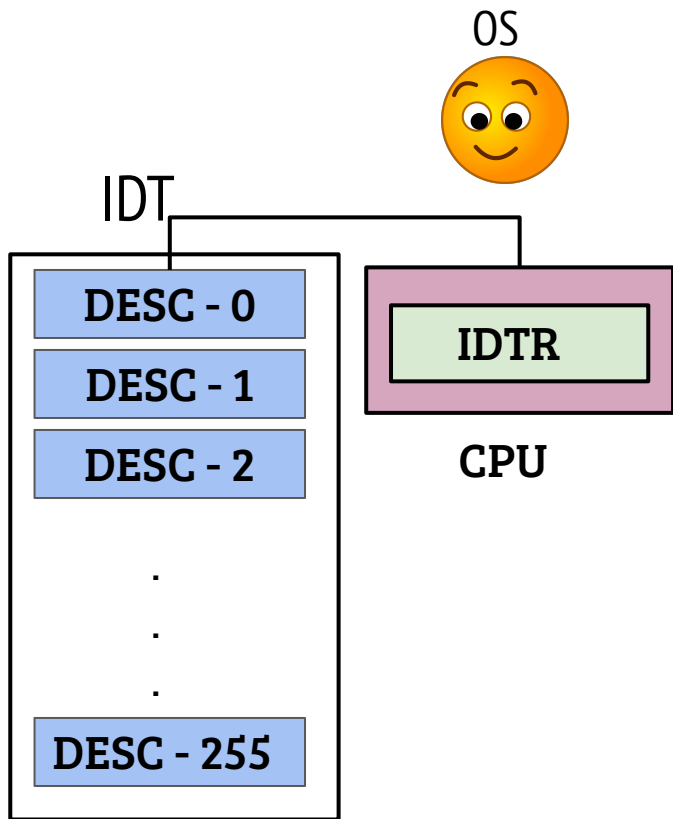
- CR3 register points to the address space translation information
- When executed from user space results in protection fault
- “mov” instruction is not privileged per se, but the operand is privileged

# Interrupt Descriptor Table (IDT): gateway to handlers



- Interrupt descriptor table provides a way to define handlers for different events like external interrupts, faults and system calls by defining the descriptors
- Descriptors 0-31 are for predefined events e.g., 0 → Div-by-zero exception etc.
- Events 32-255 are user defined, can be used for h/w and s/w interrupt handling

# Defining the descriptors (OS boot)



- Each descriptor contains information about handling the event
  - Privilege switch information
  - Handler address
- The OS defines the descriptors and loads the IDTR register with the address of the descriptor table (using *LIDT* instruction)

# System call INT instruction (Conventional Method)

- INT #N: Raise a software interrupt. CPU invokes the handler defined in the IDT descriptor #N (if registered by the OS)
- Conventionally, IDT descriptor 128 (0x80) is used to define system call entry gates
- The generic system call handler invokes the appropriate handler function.  
How?


# System call INT instruction (Conventional Method)

- INT #N: Raise a software interrupt. CPU invokes the handler defined in the IDT descriptor #N (if registered by the OS)
- Conventionally, IDT descriptor 128 (0x80) is used to define system call entry gates
- The generic system call handler invokes the appropriate handler function, How?
  - Every system call is associated with a number (defined by OS)
  - User process sends information like system call number, arguments through CPU registers which is used to invoke the actual handler

# System call in Linux Kernel (using syscall inst.)


- X86 provides a fast system call method through the “syscall” instruction
- OS configures designated privileged registers with the entry address (and other information related to privilege change)
- The hardware saves the next instruction address (user return address) into RCX, change privilege levels and sets RIP to the syscall entry address. (SP and CR3 are not modified)
- Arguments and return value
  - RAX: System call # and return value
  - Arguments passed: RDI, RSI, RDX, R10, R8, R9

# Post-boot OS execution



- OS execution is triggered because of interrupts, exceptions or system calls


# Post-boot OS execution



- OS execution is triggered because of interrupts, exceptions or system calls
- Exceptions and interrupts are abrupt, the user process may not be prepared for this event to happen. What can go wrong and how to handle it?



# Post-boot OS execution



- OS execution is triggered because of interrupts, exceptions or system calls
- Exceptions and interrupts are abrupt, the user process may not be prepared for this event to happen. What can go wrong and how to handle it?
- The interrupted program may become corrupted after resume! The OS need to save the user execution state and restore it on return

# Post-boot OS execution

External events a.k.a  
Interrupts

Kernel mode

Software interrupts  
(INT instructions)

- Does the OS need a separate stack?
- How many OS stacks are required?
- How the user process state preserved on entry to OS and restored on return to user space?
- Which address space the OS uses?

for this event to happen. What can go wrong and how to handle it?

- The interrupted program may become corrupted after resume! The OS need to save the user execution state and restore it on return

# The OS stack

- OS execution requires a stack for obvious reasons (function call & return)
- Can the OS use the user stacks?

# The OS stack

- OS execution requires a stack for obvious reasons (function call & return)
- Can the OS use the user stacks?
- No. Because of security and efficiency reasons,
  - The user may have an invalid SP at the time of entry
  - OS need to erase the used area before returning


# The OS stack

- OS execution requires a stack for obvious reasons (function call & return)
- Can the OS use the user stacks?
- No. Because of security and efficiency reasons,
  - The user may have an invalid SP at the time of entry
  - OS need to erase the used area before returning
- If OS has its own stack, who switches the stack on kernel entry?

# The OS stack

- OS execution requires a stack for obvious reasons (function call & return)
- Can the OS use the user stacks?
- No. Because of security and efficiency reasons,
  - The user may have an invalid SP at the time of entry
  - OS need to erase the used area before returning
- If OS has its own stack, who switches the stack on kernel entry?
- On X86 systems, the hardware (or OS in case of “syscall”) switches the stack pointer to the stack address configured by the OS

# Post-boot OS execution



- Does the OS need a separate stack?
- Yes, the hardware switches the SP to point it to a configured OS stack
- How many OS stacks are required?
- How the user process state preserved on entry to OS and restored on return to user space?
- Which address space the OS uses?

The interrupted program may become corrupted after resuming. The OS needs to save the user execution state and restore it on return.

# Management of OS stacks

- A per-process OS stack is required to allow multiple processes to be in OS mode of execution simultaneously
- Working?



# Management of OS stacks

- A per-process OS stack is required to allow multiple processes to be in OS mode of execution simultaneously
- Working
  - The OS configures the kernel stack address of the currently executing process in the hardware
  - The hardware switches the stack pointer on system call or exception

# Management of OS stacks

- A per-process OS stack is required to allow multiple processes to be in OS mode of execution simultaneously
- Working
  - The OS configures the kernel stack address of the currently executing process in the hardware
  - The hardware switches the stack pointer on system call or exception
- What about external interrupts?

# Management of OS stacks

- A per-process OS stack is required to allow multiple processes to be in OS mode of execution simultaneously
- Working
  - The OS configures the kernel stack address of the currently executing process in the hardware
  - The hardware switches the stack pointer on system call or exception
- What about external interrupts?
  - Separate interrupt stacks are used by OS for handling interrupts

# Post-boot OS execution

External events a.k.a  
Interrupts


Kernel mode

Software interrupts  
(INT instructions)

- Does the OS need a separate stack?
- Yes, the hardware switches the SP to point it to a configured OS stack
- How many OS stacks are required?
- For every process, a kernel stack is required
- How is the user process state preserved on entry to OS and restored on return to user space?
- Which address space the OS uses?

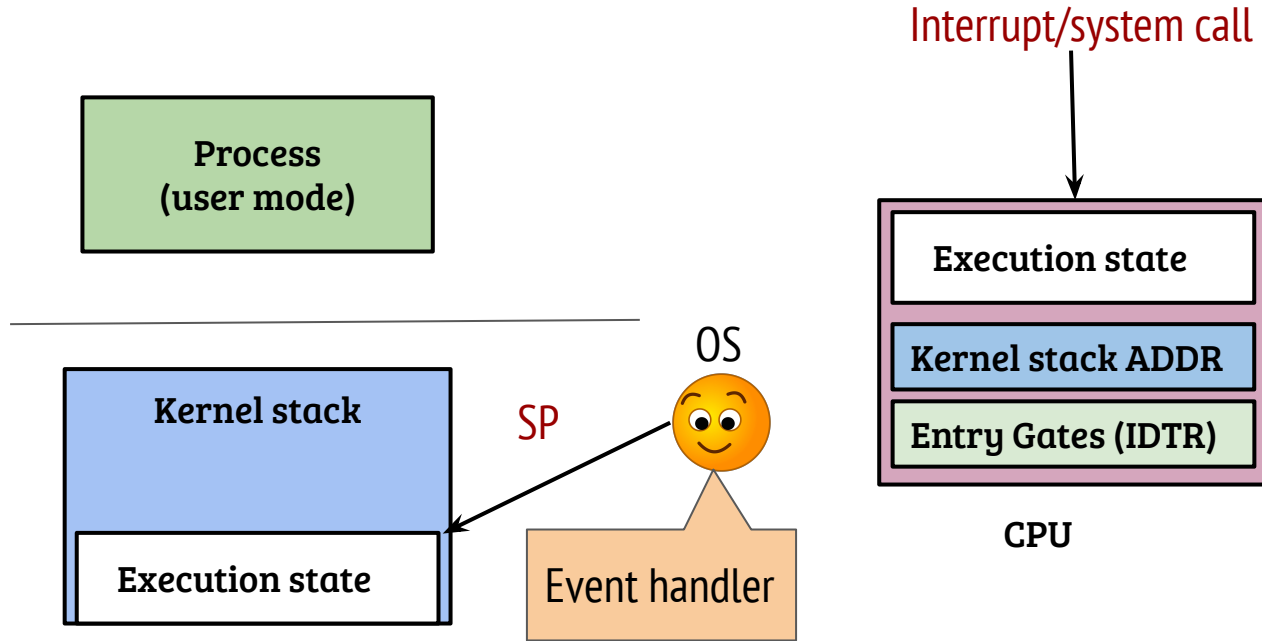
The interrupted program may become corrupted after resume. The OS needs to save the user execution state and restore it on return.

# User-kernel context switch



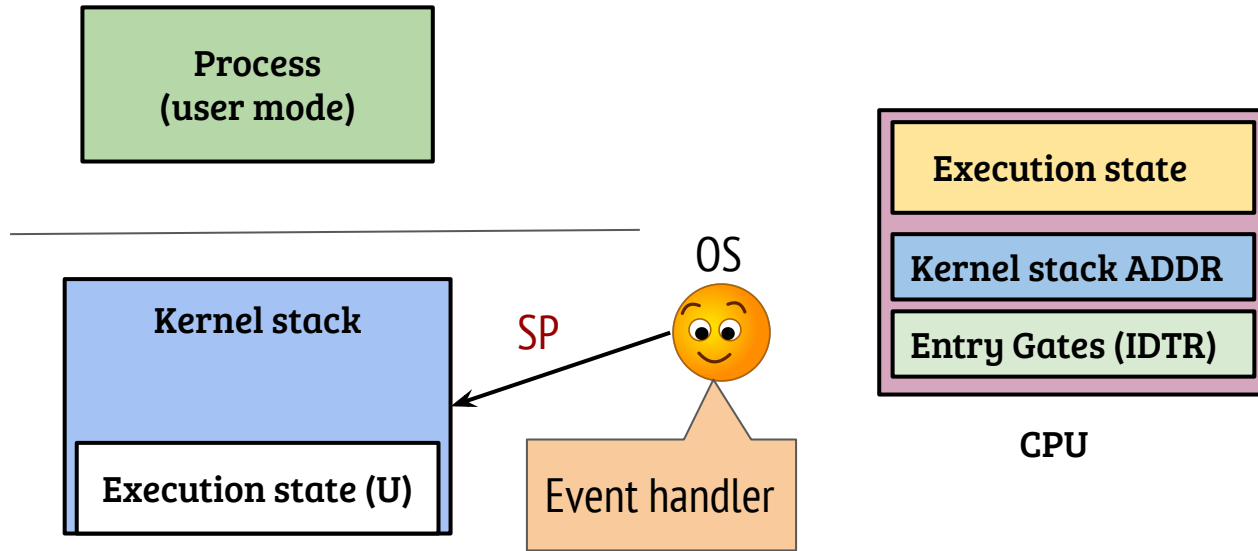
- The OS configures the kernel stack of the process before scheduling the process on the CPU

# User-kernel context switch



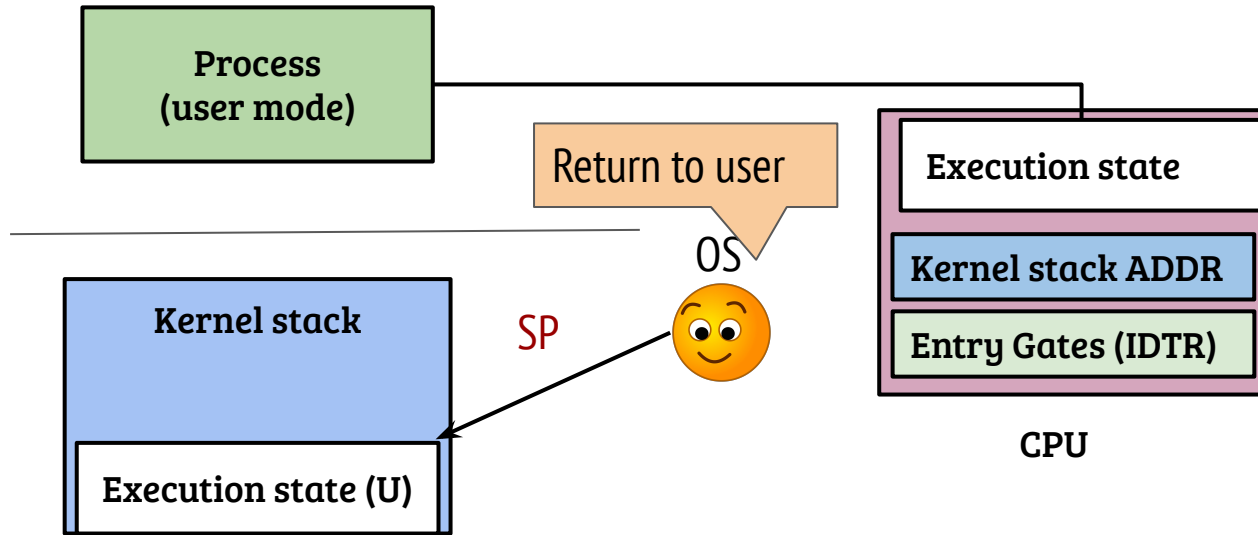
- The CPU saves the execution state onto the kernel stack
- The OS handler finds the SP switched with user state saved (fully or partially depending on architectures)

# User-kernel context switch



- The OS executes the event (syscall/interrupt) handler
  - Makes use of the kernel stack
  - Execution state on CPU is of OS at this point

# User-kernel context switch




- The kernel stack pointer should point to the position at the time of entry
- CPU loads the user execution state and resumes user execution



# Post-boot OS execution

- Does the OS need a separate stack?
- Yes, the hardware switches the SP to point it to a configured OS stack
- How many OS stacks are required?
- For every process, a kernel stack is required
- How the user process state preserved on entry to OS and restored on return to user space?
- The user execution state is saved/restored using the kernel stack by the hardware (and OS)
- Which address space the OS uses?


# The OS address space



OS

Not only I have to enable address space for each process, I need an address space myself which is protected from the user processes. Design?

# The OS address space



OS


Not only I have to enable address space for each process, I need an address space myself which is protected from the user processes. Design?

- Two possible design approaches
  - Use a separate address space for the OS, change the translation information on every OS entry (inefficient)
  - Consume a part of the address space from all processes and protect the OS addresses using H/W assistance (most commonly used)

# Post-boot OS execution

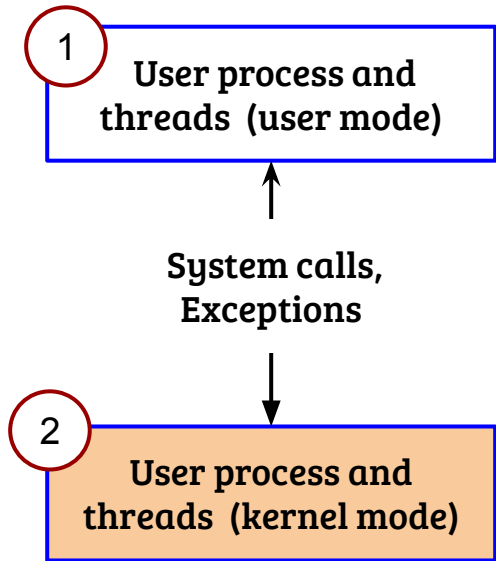
- Does the OS need a separate stack?
- Yes, the hardware switches the SP to point it to a configured OS stack
- How many OS stacks are required?
- For every process, a kernel stack is required
- How the user process state preserved on entry to OS and restored on return to user space?
- The user execution state is saved/restored using the kernel stack by the hardware (and OS)
- Which address space the OS uses?
- A part of the process address space is reserved for OS and is protected

# Execution contexts in Linux



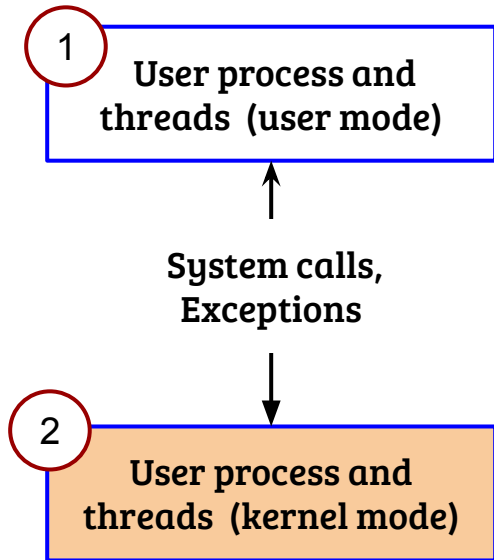
- In a linux system, the CPU can be executing in one of the above contexts
- For (3), (4) and (5), the context is not associated with any user process

# User contexts




- What are the changes in the CPU state? {CPL, Stack, CR3}
- Can a process sleep { in (1) and (2) }?
- Can a process in user mode preempted?
- Can a process in kernel mode preempted?

# User contexts



- What are the changes in the CPU state? {CPL, Stack, CR3}
- CPL and Stack change, CR3 changes if PTI enabled
- Can a process sleep { in (1) and (2) }?
- Yes, it can (lock holding conditions apply for 2)
- Can a process in user mode be preempted?
- Yes
- Can a process in kernel mode be preempted?
- Yes (if not explicitly disabled)

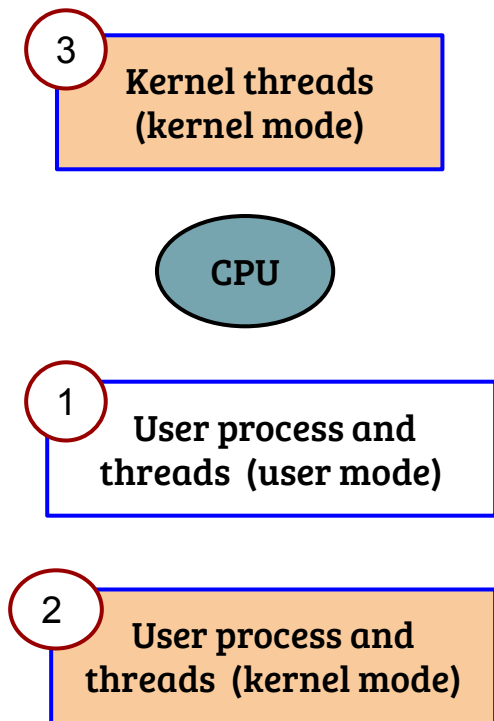
# Kernel threads



- Kernel threads are independent of user processes and threads
- Created in kernel using *kthread\_create*
- How is a kernel thread different?
- Can it sleep?
- Can it be preempted?
- Which contexts can preempt a kernel thread?



# Kernel threads



- How is a kernel thread different?
- Kernel thread never executes in user mode
- Does not require a MM context of its own
- Can it sleep?
- Yes, it can (lock holding conditions apply)
- Can it be preempted?
- Yes (if not explicitly disabled)
- Which contexts can preempt a kernel thread?
- User, Interrupt and SoftIRQ

# Hardware interrupts (Background)

5

Interrupt handler  
(kernel mode)

CPU

- Why interrupts?
- Example: Receive a packet from network
- What are the architectural support?

# Hardware interrupts (Background)

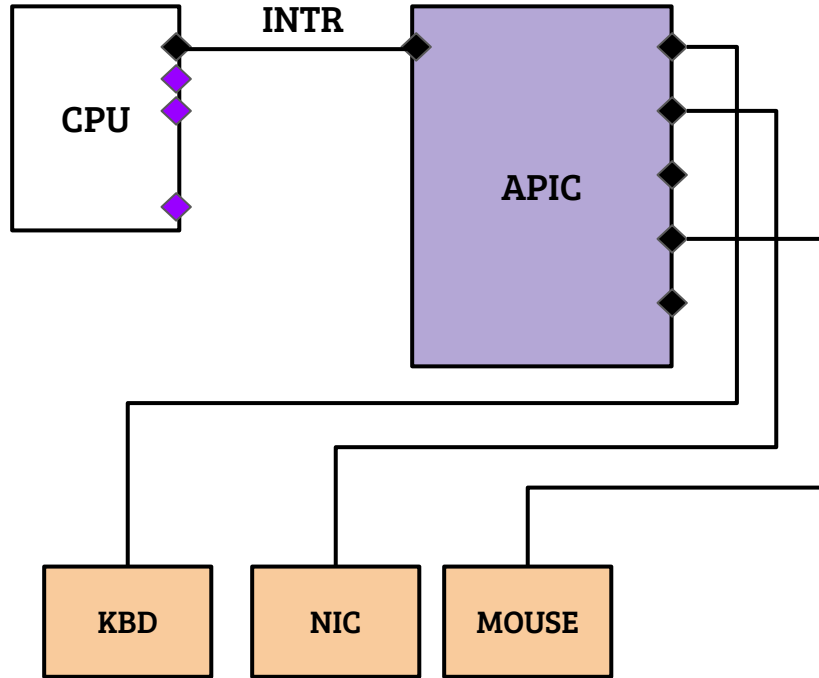
5

Interrupt handler  
(kernel mode)

CPU


- Why interrupts?
- Example: Receive a packet from network
- Avoid CPU wastage due to polling
- Responsive and scalable systems
- What are the architectural support?
- CPU has limited #of interrupt PINs → How to multiplex many devices?

# Interrupt architecture - PIC and APIC




- Every device attached to the APIC is configured with a unique IRQ number
- APIC saves the IRQ in a control port register and raise CPU interrupt line on receipt of device interrupt
- CPU reads the IRQ number and invokes the interrupt handler
- Waits for acknowledgement before clearing the INTR line
- Selective disabling of IRQs possible
  - != cli (CPU interrupt disable)
  - New interrupts not lost

# Interrupt handling




- IDT configured to load the interrupt execution context (CPL and stack)
- Interrupt entry: save regs, switch CR3 if needed
- `do_IRQ` checks the descriptor flags and invokes the real handler
- The device driver handler implements the device specific functionalities
- When is the interrupt acknowledged (i.e., INTR is cleared)?
- How long is the device interrupt masked?
- Not all interrupts can be handled quickly, e.g., NIC RCV

# Interrupt handling in three stages




- Critical tasks: Interrupt context setup, APIC acknowledgement
- Semicritical: Accessing/updating device state, e.g., update receive queue pointers of a NIC
- Deferable: Actions that are device independent e.g., Network stack processing

# Interrupt handling: SoftIRQ



- Carry out deferrable operations, can be preempted by interrupts
- Like an interrupt, it can be raised, disabled, enabled, masked
- Executed by the local CPU kernel thread (*ksoftirqd*, one per CPU)
  - Infinite loop checking for pending softIRQ (set when softirq is raised)
  - Often scheduled on `irq_exit( )` or explicit wakeup


# Interrupt context



- What are the changes in the CPU state? {CPL, Stack, CR3}
- Can an interrupt handler sleep?
- Can it be preempted?




# Interrupt context




- What are the changes in the CPU state? {CPL, Stack, CR3}
- CPL and Stack change (interrupt stack used), CR3 changes if entering from user mode in a PTI enabled system
- Can an interrupt handler sleep?
- No, Linux does not allow sleeping (directly/indirectly) in an interrupt handler
- Can it be preempted?
- Only by another interrupt (if APIC Acked and interrupts enabled on CPU )

# Contexts in action: network receive

User process



Kernel thread  
(ksoftirqd)




- The user process invokes `recv( )` system call (blocking)
- No processed payload found, the process is descheduled and put into a wait queue
- `ksoftirqd` is either suspended or processing other pending `softIRQs`

NIC


# Contexts in action: network receive

User process




Wait

Kernel thread  
(ksoftirqd)



Interrupt handler




Packet

- The NIC copies the packet (using DMA) into memory buffers (a.k.a. skbuffs) and triggers the interrupt
- Before the device specific interrupt handling, APIC is acknowledged
- The device interrupt handler update the device state while masking device interrupts
- Queues the packet for further processing and triggers a softIRQ


# Contexts in action: network receive

User process




Wait

Kernel thread  
(ksoftirqd)



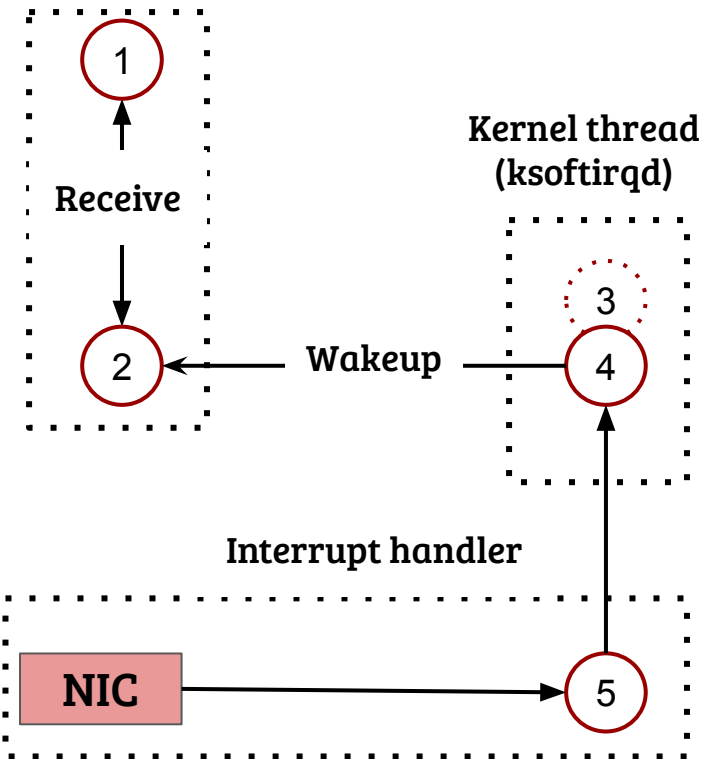
Interrupt handler



- The softIRQ is scheduled using the ksoftirqd kernel thread context
- Protocol stack processing is performed in this context
- As part of the protocol processing, the destination process is derived

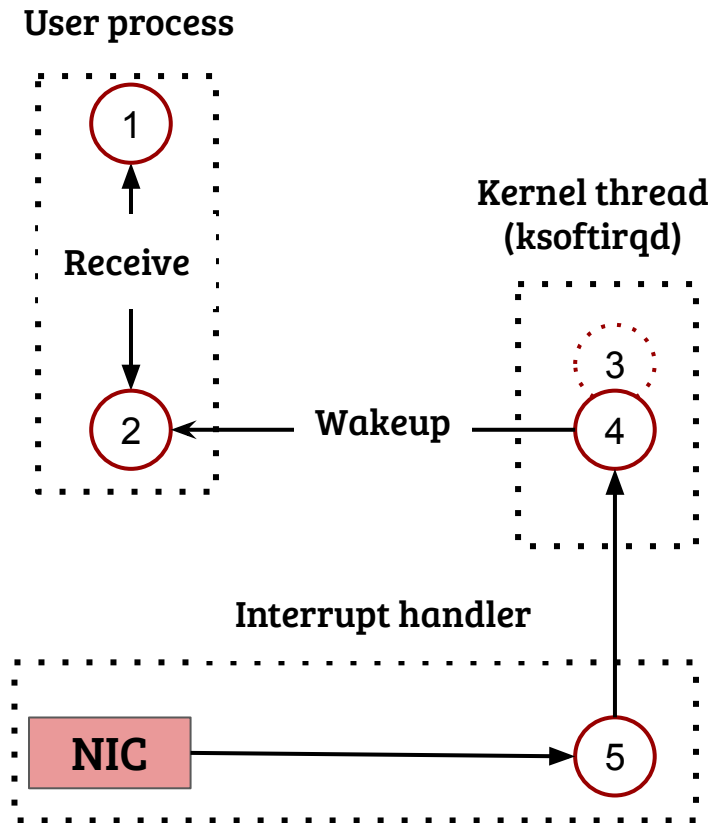
# Contexts in action: network receive

User process



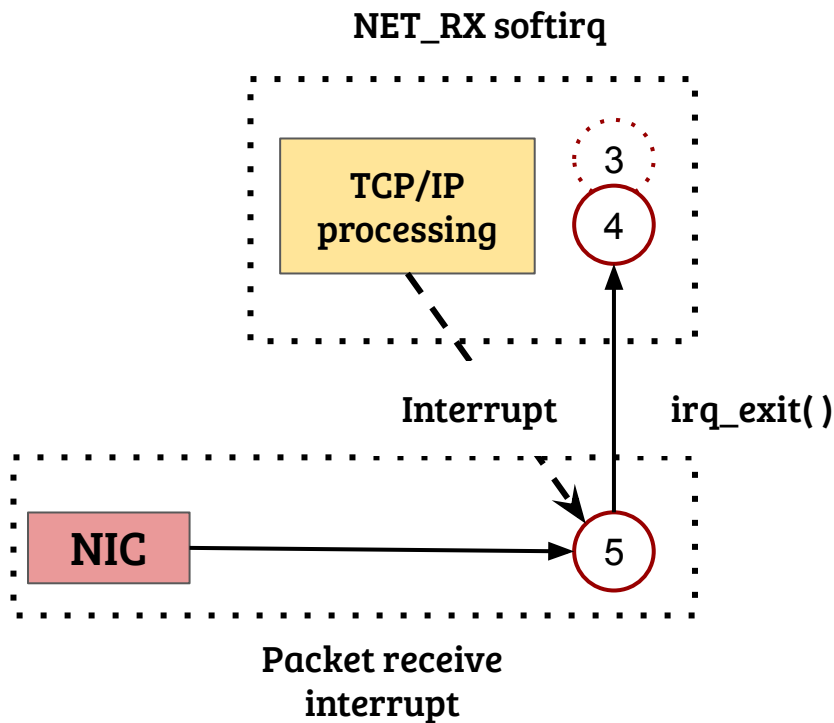
- The softIRQ processing wakes up the user process
- The user process returns from syscall (copy payload to user)
- Now, what could be the issues with this approach?

# Challenges in network receive



- Minimize network packet copy across the contexts
- Precise scheduling: application progress and fairness
- Network is always overdriven and self-adjusting in nature → rate limit as early as possible
- Issues
  - Receive livelock: CPU is always handling interrupts
  - User process starvation due to softIRQ processing

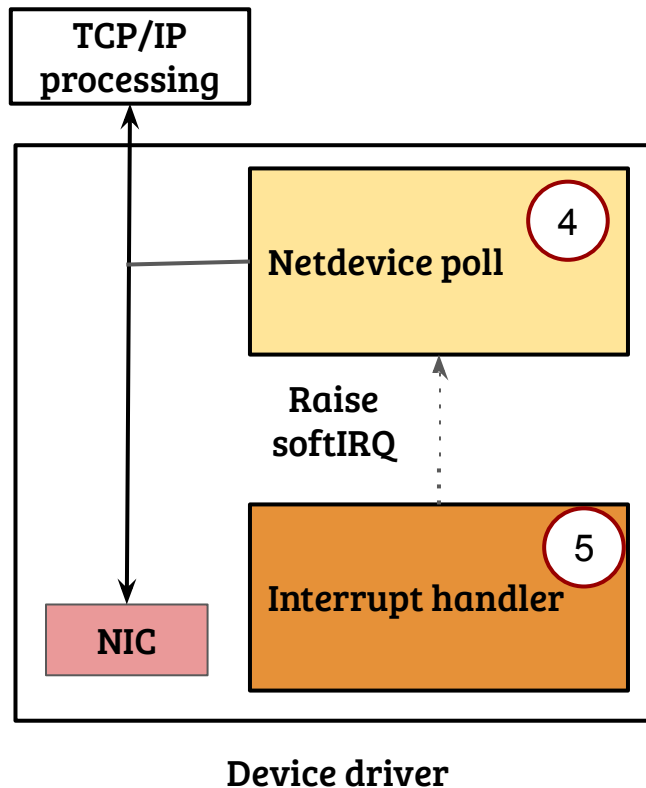
# Receive livelock <sup>1</sup>



- Root cause: Interrupts have the highest priority over other contexts
- If the rate of interrupts is high, the system remains in interrupt handling mode, resulting in *receive livelock*
- Solution approach: Lower the priority of interrupts under heavy load
- How?

1. <https://www.usenix.org/legacy/publications/library/proceedings/sd96/mogul.html>

# NAPI: Interrupt + Polling



- Interrupt handler raises softIRQ after disabling packet receive interrupts
- Driver registered poll method is invoked
  - Executes till receive queue is empty or an upper threshold (budget)
  - Enable the interrupt (if queue is empty) and return
- Advantages
  - Low network load, more interrupt driven
  - High load, less interrupt processing
  - Avoid wasted work, drop packets early (in the device buffer)



# Context related helper routines

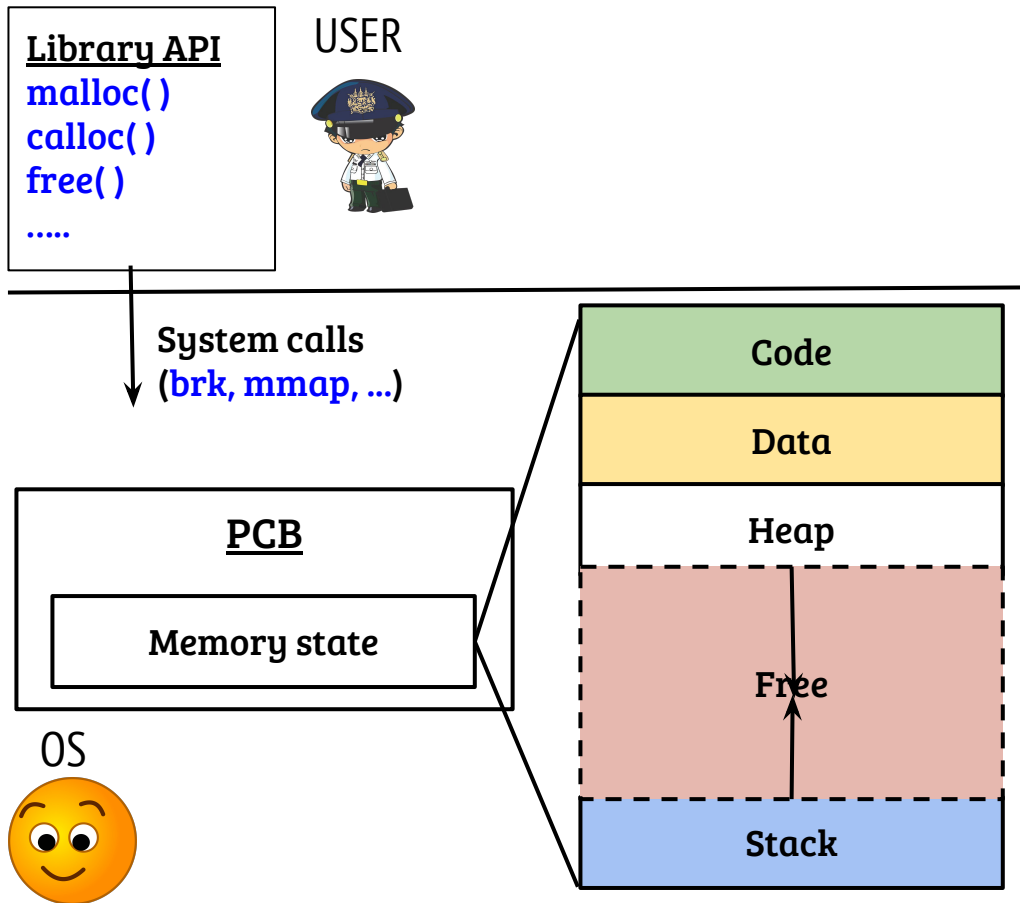
- `bool in_irq( )`
  - True if the current execution is in hardware interrupt
- `bool in_softirq( )`
  - True if the current execution is in a softIRQ or it is disabled
- `bool in_interrupt( )`
  - True if we are in NMI, IRQ, softIRQ context or have softIRQs disabled
- `bool in_task( )`
  - True if executing in a task context, *current* is valid
- Disabling/enabling interrupts
  - `local_irq_disable/enable( )`
- Disabling/enabling softIRQs
  - `local_bh_disable/enable( )`

# CS614: Linux Kernel Programming

## Virtual Memory

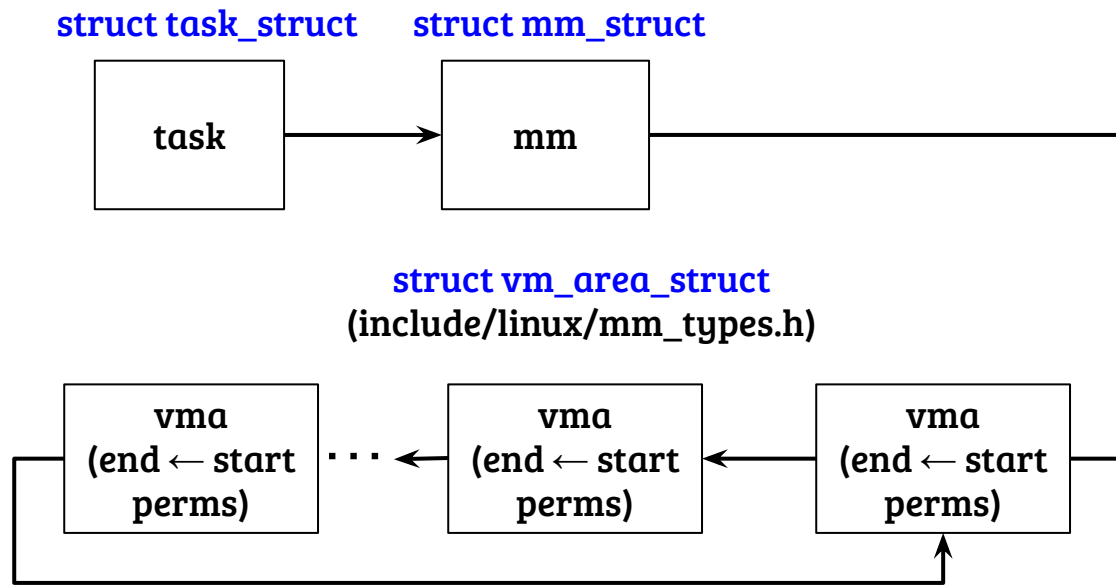
Debadatta Mishra, CSE, IIT Kanpur

# User API for memory management



- Generally, user programs use library routines to allocate/deallocate memory
- OS provides some address space manipulation system calls (today's agenda)

# Virtual memory management




- start and end never overlaps between two vm areas
- can merge/extend vmas if permissions match
- linux maintains both `rb_tree` and a sorted list (see `mm/filemap.c`)

The OS implements VM system calls like `mmap( )`, `mprotect( )` by manipulating the VMAs

# Address translation: Paging

- The idea of paging
  - Partition the address space into fixed sized blocks (call it pages)
  - Physical memory partitioned in a similar way (call it page frames)
  - OS creates a mapping between *page* to *page frame* , H/W uses the mapping to translate VA to PA
- With increased address space size, single level page table entry is not feasible, because
  - Increasing page size increases internal fragmentation
  - Small pages may not be suitable to hold all mapping entries

# 4-level page tables: 48-bit VA (Intel x86\_64)



- Virtual address size =  $2^{48}$ , Page size = 4096 bytes
- Four-levels of page table, entry size = 64 bits

# Paging example (structure of an example PTE)



- PFN occupies a significant portion of PTE entry (8 bits in this example)

P

Present bit, 1  $\Rightarrow$  entry is valid

W

Write bit, 1  $\Rightarrow$  Write allowed

S

Privilege bit, 0  $\Rightarrow$  only kernel mode access is allowed

A

Accessed bit, 1  $\Rightarrow$  Address accessed (set by H/W during walk)

D


Dirty bit, 1  $\Rightarrow$  Address written (set by H/W during walk)

X

Execute bit, 1  $\Rightarrow$  Instruction fetch allowed for this page

Reserved/unused bits

# 4-level page tables: example translation



Virtual address = 0x180001008

- Hardware translation by repeated access of page table stored in physical memory
- Page table entry: 12 bits LSB is used for access flags



# Paging: translation efficiency

	0x20100: mov \$0, %rax;	
	0x20102: mov %rax, (%rbp);	// sum=0
sum = 0;	0x20104: mov \$0, %rcx;	// ctr=0
for(ctr=0; ctr<10; ++ctr)	0x20106: cmp \$10, %rcx;	// ctr < 10
sum += ctr;	0x20109: jge 0x2011f;	// jump if >=
	0x2010f: add %rcx, %rax;	
	0x20111: mov %rax, (%rbp);	// sum += ctr
	0x20113: inc %rcx	// ++ctr
	0x20115: jmp 0x20106	// loop
	0x2011f: .....	

- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

# Paging: translation efficiency

```
0x20100: mov $0, %rax;
```

```
0x20102: mov %rax, (%rbp); // sum=0
```

- Instruction execution: Loop =  $10 * 6$ , Others =  $2 + 3$ 
    - Memory accesses during translation =  $65 * 4 = 260$
  - Data/stack access: Initialization = 1, Loop = 10
    - Memory accesses during translation =  $11 * 4 = 44$
  - A lot of memory accesses ( $> 300$ ) for address translation
  - How many distinct pages are translated?
- 
- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

# Paging with TLB: translation efficiency

Translate(V){

PageAddress P = V >> 12;

TLBEntry entry = lookup(P);

if (entry.valid) return entry.pte;

entry = PageTableWalk(V);

MakeEntry(entry);

return entry.pte;

}

TLB	
Page	PTE
0x20	0x750
0x7FFF	0x890

- TLB is a hardware cache which stores *Page* to *PFN* mapping
- After first miss for instruction fetch address, all others result in a TLB hit
- Similarly, considering the stack virtual address range as 0x7FFF000 - 0x8000000, one entry in TLB avoids page table walk after first miss


# Paging: translation efficiency

```
0x20100: mov $0, %rax;
```

- Instruction execution: Loop =  $10 * 6$ , Others =  $2 + 3$ 
  - Memory accesses during translation =  $65 * 4 = 260$
- Data/stack access: Initialization = 1, Loop = 10
  - Memory accesses during translation =  $11 * 4 = 44$
- A lot of memory accesses ( $> 300$ ) for address translation
- How many distinct pages are translated?
- One code page (0x20) and one stack page (0x7FFF). Caching these translations, will save a lot of memory accesses.


required (for translation) during the execution of the above code?

# Address translation (TLB + PTW)




- TLB in the path of address translation
- Separate TLBs for instruction and data, multi-level TLBs
- In X86, OS can not make entries into the TLB directly, it can flush entries

# Address translation (TLB + PTW)



- TLB in the path of address

- How TLB is shared across multiple processes?
- Why page fault is necessary?
- How OS handles the page fault?



into the TLB directly, it can flush entries

# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution

# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution
- Process B may be using the same addresses used by A. Result: Wrong translation



# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution
- Correctness ensured. Performance is an issue (with frequent context switching)

# TLB: Sharing across applications

Process (A)

Process (B)

- Assume that, process A is currently executing. What happens when process B is scheduled?

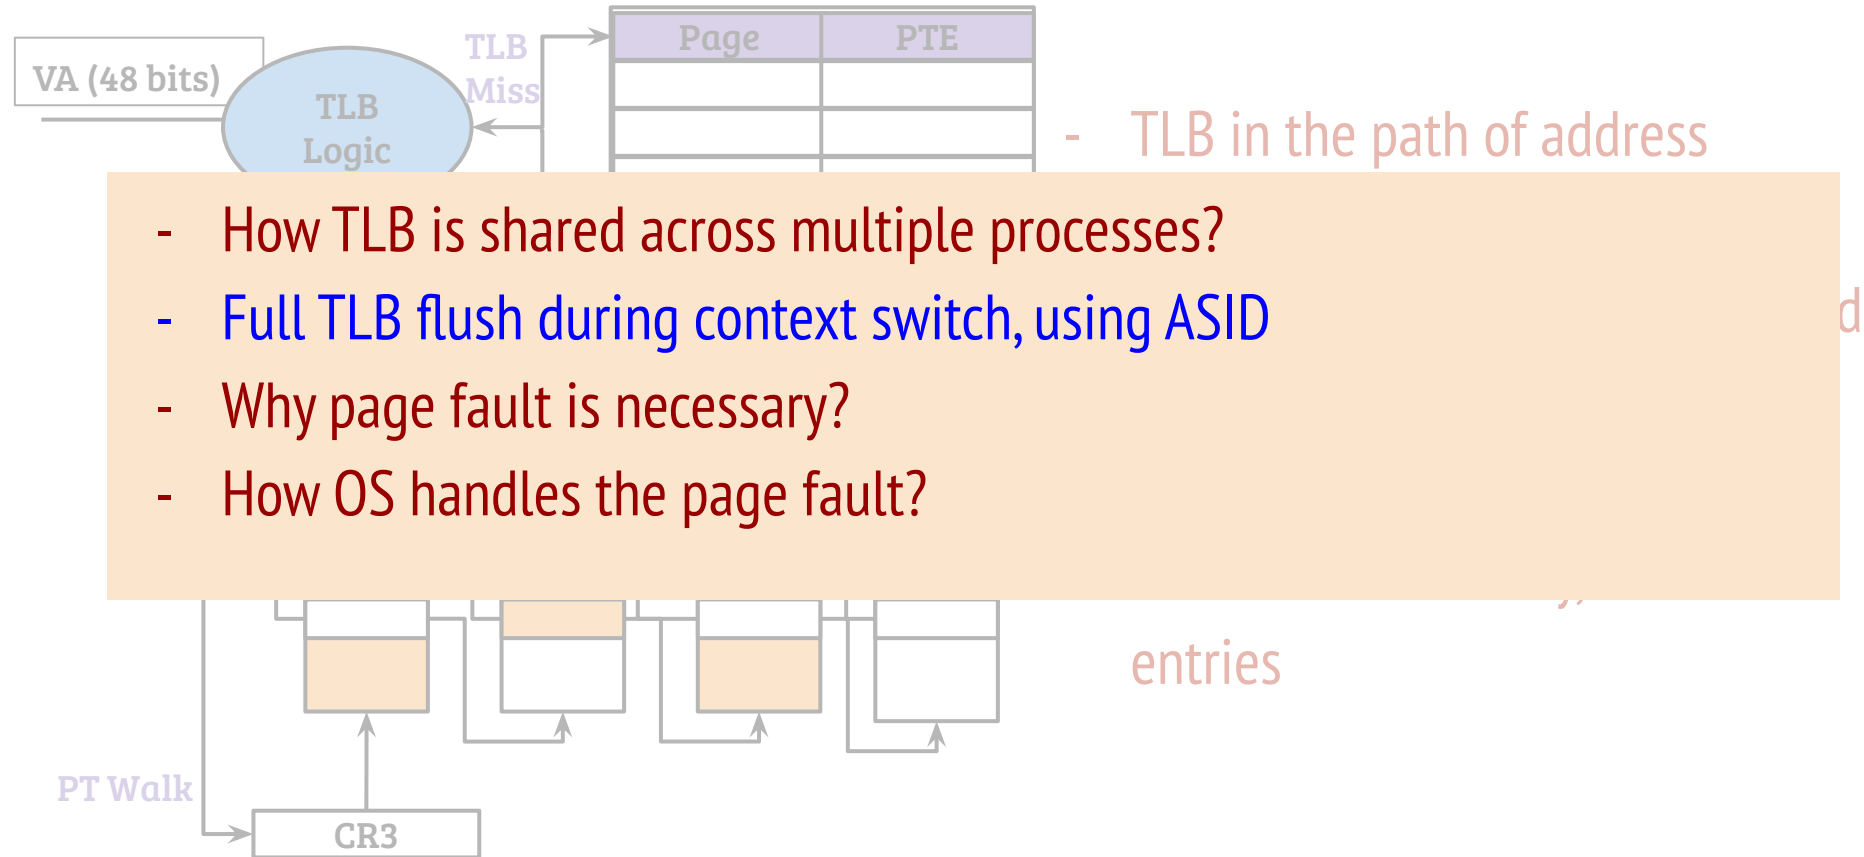
- A) Do nothing
- B) Flush the whole TLB
- C) Some other solution

- Address space identified (ASID) along with each TLB entry to identify the process

ASID	Page	PTE
A	0x100	0x200007
A	0x101	0x205007
B	0x100	0x301007
B	0x101	0x302007

TLB


# Address translation (TLB + PTW)



# Address translation (TLB + PTW)



- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?



# Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

# Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Error code

Other and unused	I	R	U	W	P
------------------	---	---	---	---	---

P

**Present bit, 1  $\Rightarrow$  fault is due to protection**

W

**Write bit, 1  $\Rightarrow$  Access is write**

U

**Privilege bit, 1  $\Rightarrow$  Access is from user mode**

R

**Reserved bit, 1  $\Rightarrow$  Reserved bit violation**

I

**Fetch bit, 1  $\Rightarrow$  Access is Instruction Fetch**

- Error code is pushed into the kernel stack by the hardware

# Page fault handling in X86: OS fault handler

```
HandlePageFault( u64 address, u64 error_code)
{
    If( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn( );
        install_pte(address, PFN);
        return;
    }
    RaiseSignal(SIGSEGV);
}
```


# Address translation (TLB + PTW)

- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?
- The hardware invokes the page fault handler by placing the error code and virtual address. The OS handles the page fault either fixing it or raising a SEGFAULT.



# Swapping (swap-out)

DRAM



Swap (Hard disk)


Number of free PFNs are very few in the system. I can not break my promise made to the applications. Let me swap-out some memory. But which one to swap-out?



OS

AllocatePFN()

# Swapping (swap-out)



My page replacement policy will help me deciding the victims (V). Can I just swap-out? What if the swapped-out pages are accessed? I should be prepared for that too!




OS

AllocatePFN()

# Swapping (swap-out)

DRAM



PTE mapping the victim PFN (before swap)



PTE mapping the victim PFN (after swap)



Swap (Hard disk)

Update the present-bit to 0 in the PTE such that any access to the page through the virtual address will result in a page fault. Also maintain the swap address in the PTE.

OS



AllocatePFN()

# Swapping (swap-out)


DRAM



**PTE mapping the victim PFN (before swap)**



**PTE mapping the victim PFN (after swap)**



Swap (Hard disk)

Content of the PFN is now in the swap device. In future, any translation using the PTE will result in a page fault. The page fault handler would copy it back from the swap device.



OS

AllocatePFN()

# Page fault with swap-in

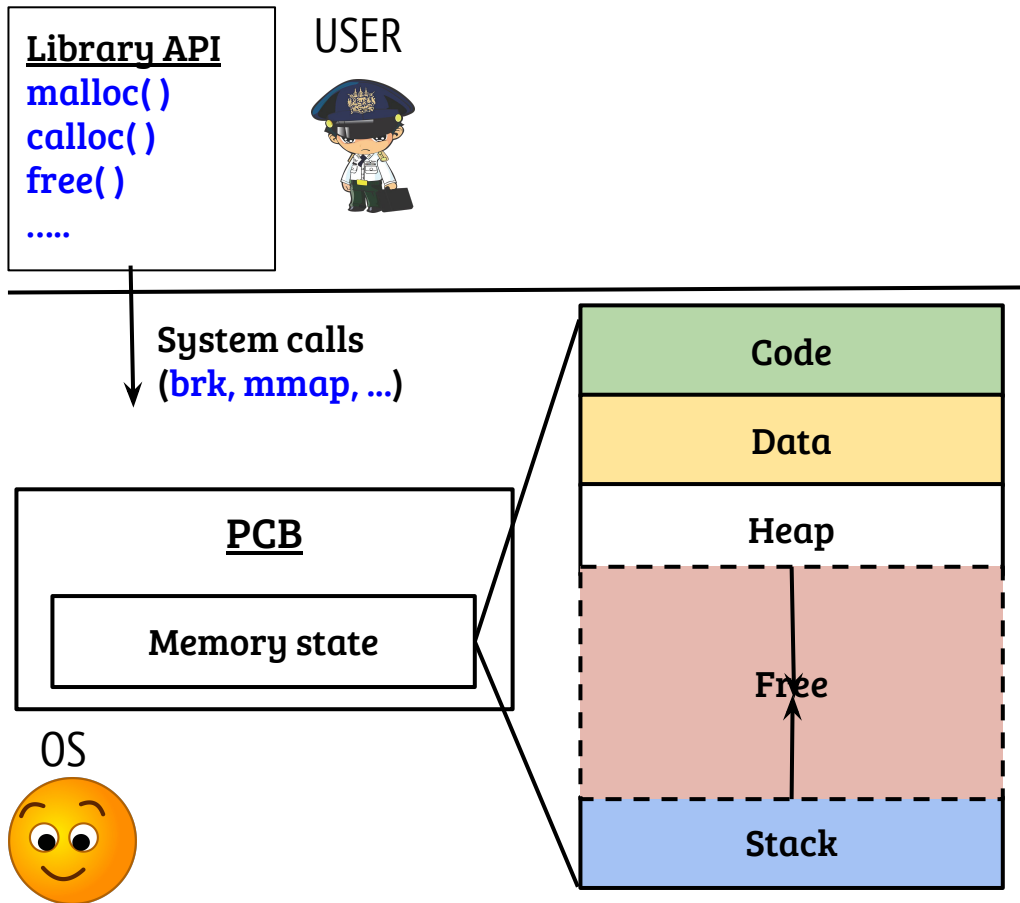
```
HandlePageFault( u64 address, u64 error_code)
{
    If ( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn();
        If ( is_swapped_pte(address) )    // Check if the PTE is swapped out
            swapin(getPTE(address), PFN); // Copy the swap block to PFN
        install_pte(address, PFN);      // and update the PTE
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

# CS614: Linux Kernel Programming

## Virtual Memory


Debadatta Mishra, CSE, IIT Kanpur

# User API for memory management



- Generally, user programs use library routines to allocate/deallocate memory
- OS provides some address space manipulation system calls (today's agenda)

# Virtual memory management



- start and end never overlaps between two vm areas
- can merge/extend vmas if permissions match
- linux maintains both `rb_tree` and a sorted list (see `mm/filemap.c`)


The OS implements VM system calls like `mmap( )`, `mprotect( )` by manipulating the VMAs



# Address translation: Paging

- The idea of paging
  - Partition the address space into fixed sized blocks (call it pages)
  - Physical memory partitioned in a similar way (call it page frames)
  - OS creates a mapping between *page* to *page frame* , H/W uses the mapping to translate VA to PA
- With increased address space size, single level page table entry is not feasible, because
  - Increasing page size increases internal fragmentation
  - Small pages may not be suitable to hold all mapping entries

# 4-level page tables: 48-bit VA (Intel x86\_64)



- Virtual address size =  $2^{48}$ , Page size = 4096 bytes
- Four-levels of page table, entry size = 64 bits

# Paging example (structure of an example PTE)



- PFN occupies a significant portion of PTE entry (8 bits in this example)

P

Present bit, 1  $\Rightarrow$  entry is valid

W

Write bit, 1  $\Rightarrow$  Write allowed

S

Privilege bit, 0  $\Rightarrow$  only kernel mode access is allowed

A

Accessed bit, 1  $\Rightarrow$  Address accessed (set by H/W during walk)

D

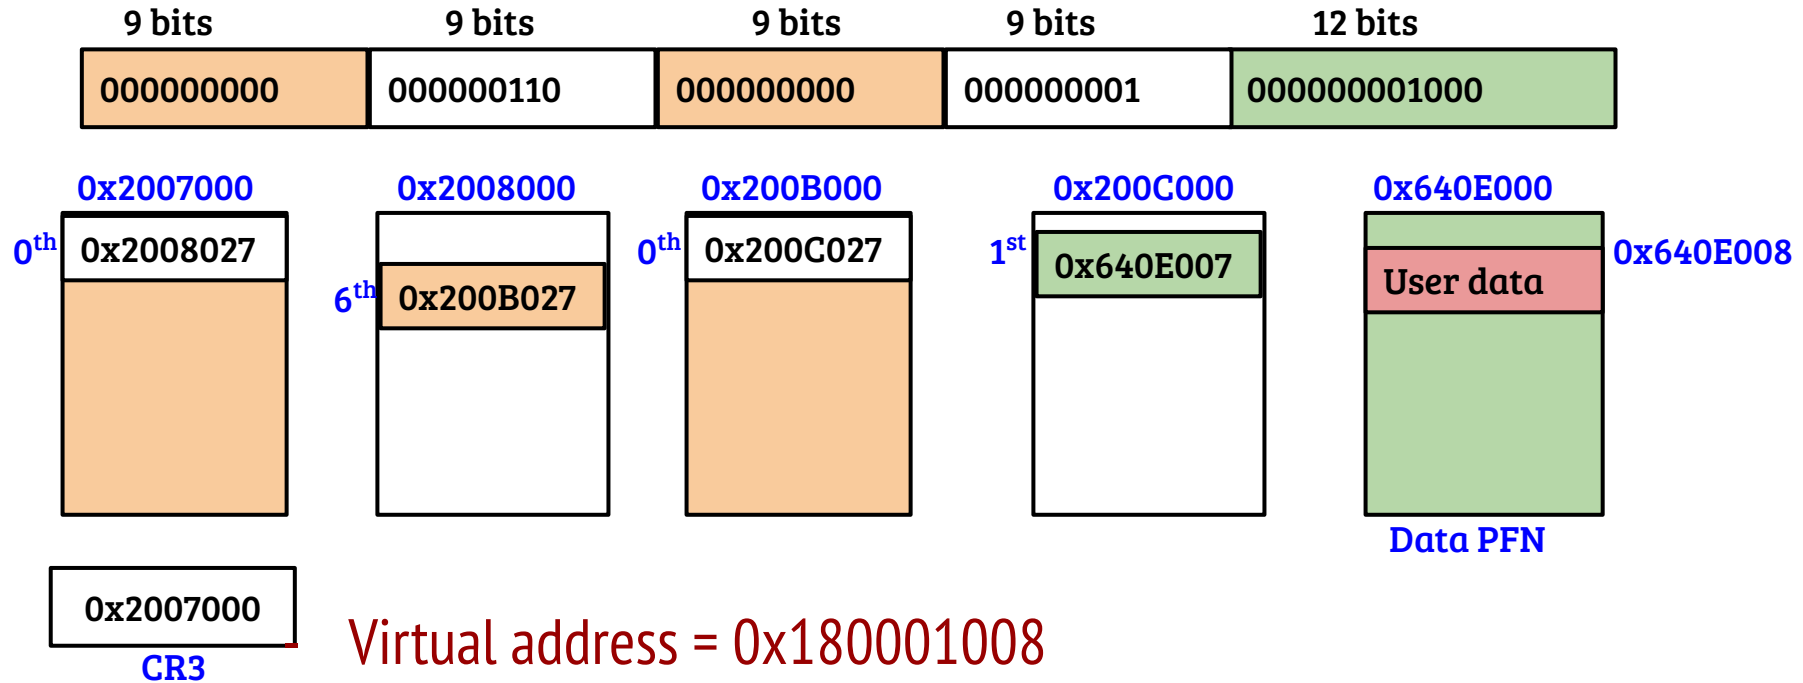
Dirty bit, 1  $\Rightarrow$  Address written (set by H/W during walk)

X

Execute bit, 1  $\Rightarrow$  Instruction fetch allowed for this page

Reserved/unused bits

# 4-level page tables: example translation



Virtual address = 0x180001008

- Hardware translation by repeated access of page table stored in physical memory
- Page table entry: 12 bits LSB is used for access flags

# Paging: translation efficiency

	0x20100: mov \$0, %rax;	
	0x20102: mov %rax, (%rbp);	// sum=0
sum = 0;	0x20104: mov \$0, %rcx;	// ctr=0
for(ctr=0; ctr<10; ++ctr)	0x20106: cmp \$10, %rcx;	// ctr < 10
sum += ctr;	0x20109: jge 0x2011f;	// jump if >=
	0x2010f: add %rcx, %rax;	
	0x20111: mov %rax, (%rbp);	// sum += ctr
	0x20113: inc %rcx	// ++ctr
	0x20115: jmp 0x20106	// loop
	0x2011f: .....	

- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

# Paging: translation efficiency

```
0x20100: mov $0, %rax;
```

```
0x20102: mov %rax, (%rbp); // sum=0
```

- Instruction execution: Loop =  $10 * 6$ , Others =  $2 + 3$ 
    - Memory accesses during translation =  $65 * 4 = 260$
  - Data/stack access: Initialization = 1, Loop = 10
    - Memory accesses during translation =  $11 * 4 = 44$
  - A lot of memory accesses ( $> 300$ ) for address translation
  - How many distinct pages are translated?
- 
- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

# Paging with TLB: translation efficiency

Translate(V){

PageAddress P = V >> 12;

TLBEntry entry = lookup(P);

if (entry.valid) return entry.pte;

entry = PageTableWalk(V);

MakeEntry(entry);

return entry.pte;

}

TLB	
Page	PTE
0x20	0x750
0x7FFF	0x890

- TLB is a hardware cache which stores *Page* to *PFN* mapping
- After first miss for instruction fetch address, all others result in a TLB hit
- Similarly, considering the stack virtual address range as 0x7FFF000 - 0x8000000, one entry in TLB avoids page table walk after first miss

# Paging: translation efficiency

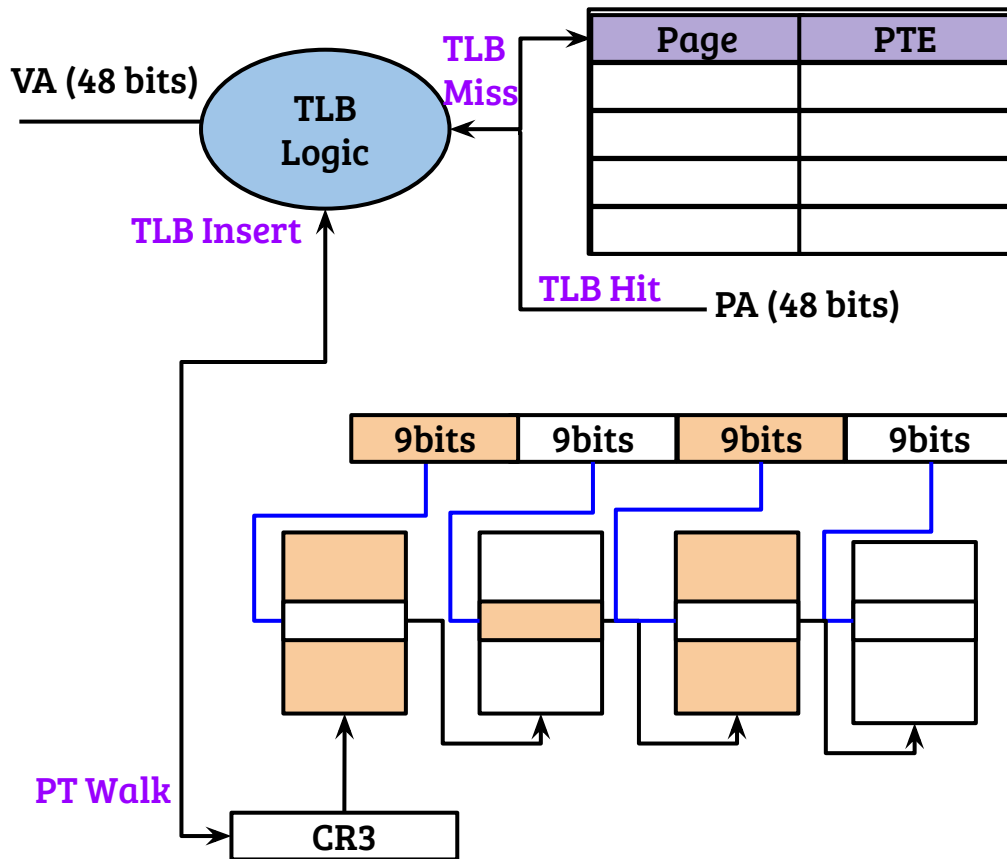
```
0x20100: mov $0, %rax;
```

- Instruction execution: Loop =  $10 * 6$ , Others =  $2 + 3$ 
  - Memory accesses during translation =  $65 * 4 = 260$
- Data/stack access: Initialization = 1, Loop = 10
  - Memory accesses during translation =  $11 * 4 = 44$
- A lot of memory accesses ( $> 300$ ) for address translation
- How many distinct pages are translated?
- One code page (0x20) and one stack page (0x7FFF). Caching these translations, will save a lot of memory accesses.

required (for translation) during the execution of the above code?




# Address translation (TLB + PTW)




- TLB in the path of address translation
- Separate TLBs for instruction and data, multi-level TLBs
- In X86, OS can not make entries into the TLB directly, it can flush entries

# Address translation (TLB + PTW)



- TLB in the path of address

- How TLB is shared across multiple processes?
- Why page fault is necessary?
- How OS handles the page fault?



into the TLB directly, it can flush entries

# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution

# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution
- Process B may be using the same addresses used by A. Result: Wrong translation

# TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
  - A) Do nothing
  - B) Flush the whole TLB
  - C) Some other solution
- Correctness ensured. Performance is an issue (with frequent context switching)

# TLB: Sharing across applications

Process (A)

Process (B)

- Assume that, process A is currently executing. What happens when process B is scheduled?

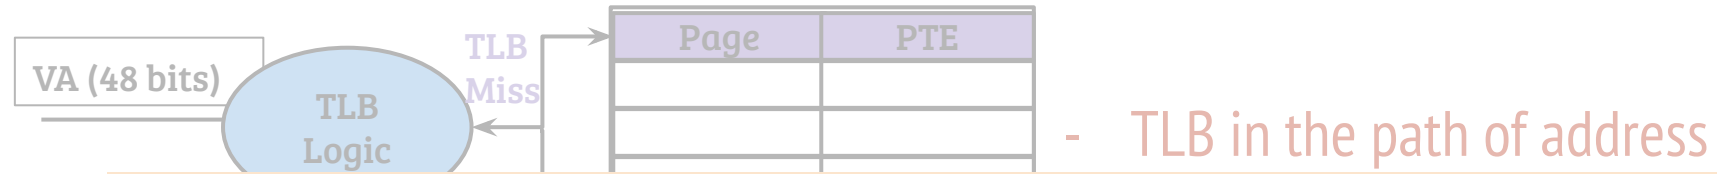
- A) Do nothing
- B) Flush the whole TLB
- C) Some other solution

- Address space identified (ASID) along with each TLB entry to identify the process


ASID	Page	PTE
A	0x100	0x200007
A	0x101	0x205007
B	0x100	0x301007
B	0x101	0x302007

TLB

# Address translation (TLB + PTW)



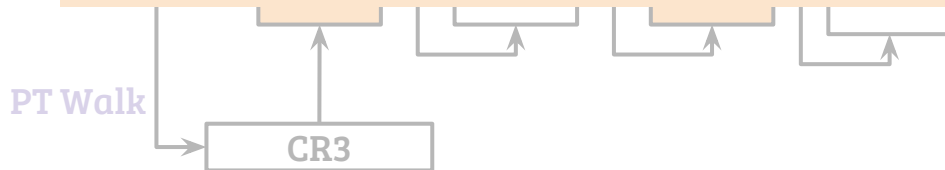
- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- How OS handles the page fault?



# Address translation (TLB + PTW)



- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?





# Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

# Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Error code

Other and unused	I	R	U	W	P
------------------	---	---	---	---	---

P

**Present bit, 1  $\Rightarrow$  fault is due to protection**

W

**Write bit, 1  $\Rightarrow$  Access is write**

U

**Privilege bit, 1  $\Rightarrow$  Access is from user mode**

R

**Reserved bit, 1  $\Rightarrow$  Reserved bit violation**

I

**Fetch bit, 1  $\Rightarrow$  Access is Instruction Fetch**

- Error code is pushed into the kernel stack by the hardware

# Page fault handling in X86: OS fault handler


```
HandlePageFault( u64 address, u64 error_code)
{
    If( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn( );
        install_pte(address, PFN);
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

# Address translation (TLB + PTW)

- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?
- The hardware invokes the page fault handler by placing the error code and virtual address. The OS handles the page fault either fixing it or raising a SEGFAULT.

# Swapping (swap-out)

DRAM



Swap (Hard disk)


Number of free PFNs are very few in the system. I can not break my promise made to the applications. Let me swap-out some memory. But which one to swap-out?



OS

AllocatePFN()

# Swapping (swap-out)



My page replacement policy will help me deciding the victims (V). Can I just swap-out? What if the swapped-out pages are accessed? I should be prepared for that too!




OS

AllocatePFN()

# Swapping (swap-out)


DRAM



PTE mapping the victim PFN (before swap)



PTE mapping the victim PFN (after swap)



Swap (Hard disk)

Update the present-bit to 0 in the PTE such that any access to the page through the virtual address will result in a page fault. Also maintain the swap address in the PTE.



OS

AllocatePFN()

# Swapping (swap-out)


DRAM



**PTE mapping the victim PFN (before swap)**



**PTE mapping the victim PFN (after swap)**



Swap (Hard disk)

Content of the PFN is now in the swap device. In future, any translation using the PTE will result in a page fault. The page fault handler would copy it back from the swap device.



OS


AllocatePFN()



# Page fault with swap-in


```
HandlePageFault( u64 address, u64 error_code)
{
    If ( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn();
        If ( is_swapped_pte(address) )    // Check if the PTE is swapped out
            swapin(getPTE(address), PFN); // Copy the swap block to PFN
        install_pte(address, PFN);      // and update the PTE
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

# Efficient translation: Huge page support




- Virtual address size =  $2^{48}$ , Page size = 2M = PFN size
- Advantages?
- Disadvantages?

# Efficient translation: Huge page support




- Virtual address size =  $2^{48}$ , Page size = 2M
- Advantages? Efficient (walk and TLB coverage)
- Disadvantages? Inefficient management

# Mixed page size support



```
walk_pmd(pmd, vaddr) {  
    if(pmd.H)  
        paddr = pmd.nextL() + (vaddr & pmask);  
    else  
        pte = pmd.nextL() + pte_offset(vaddr)  
    } // Simplified H/W logic
```

# Mixed page size support



```
walk_pmd(pmd, vaddr) {  
    if(pmd.H)  
        paddr = pmd.nextL() + (vaddr & pmask);  
    else  
        pte = pmd.nextL() + pte_offset(vaddr)  
} // Simplified H/W logic
```

- The OS may use the hardware support to implement any policy
- Transparent hugepage (THP) in Linux tries to create huge page mapping in w/o explicit user space assistance
- Policy knobs through sysfs

# Kernel Virtual Memory

- Why not treat kernel as an isolated MM context?

# Kernel Virtual Memory

- Why not treat kernel as an isolated MM context?
  - Require MM context loading/unloading on user-kernel context switch
  - In kernel context, user data is accessed (a lot!) why?
  - Even worse, user data of many processes accessed
  - In X86, a small part of the kernel can not be isolated as HW does not perform MM context switch
- Requirement: efficient memory isolation between user and kernel

# Kernel Virtual Memory


- Why not treat kernel as an isolated MM context?
  - Require MM context loading/unloading on user-kernel context switch
  - In kernel context, user data is accessed (a lot!) why?
  - Even worse, user data of many processes can be accessed
  - In X86, a small part of the kernel can not be isolated as HW does not perform MM context switch
- Requirement: efficient memory isolation between user and kernel
  - Let kernel use the same MM context of the user process
  - No context switch, no problems of accessing user data



# Kernel Virtual Memory

- Why not treat kernel as an isolated MM context?
  - Require MM context loading/unloading on user-kernel context switch
  - In kernel context, user data is accessed (a lot!) why?
  - Even worse, user data of many processes can be accessed
  - In X86, a small part of the kernel can not be isolated as HW does not perform MM context switch
- Requirement: efficient memory isolation between user and kernel
  - Let kernel use the same MM context of the user process
  - No context switch, no problems of accessing user data
- How kernel VM change propagated across processes? Isolation issues?

# Issue of Kernel VM propagation



- Kernel virtual address mapping should be present in both process page tables.
- Ex: If kernel allocates memory while serving syscall from process-1, process-2 in kernel mode should see it!
- Solution should consider that, “processes and memory are dynamically created and destroyed”

# Linux strives on family values!

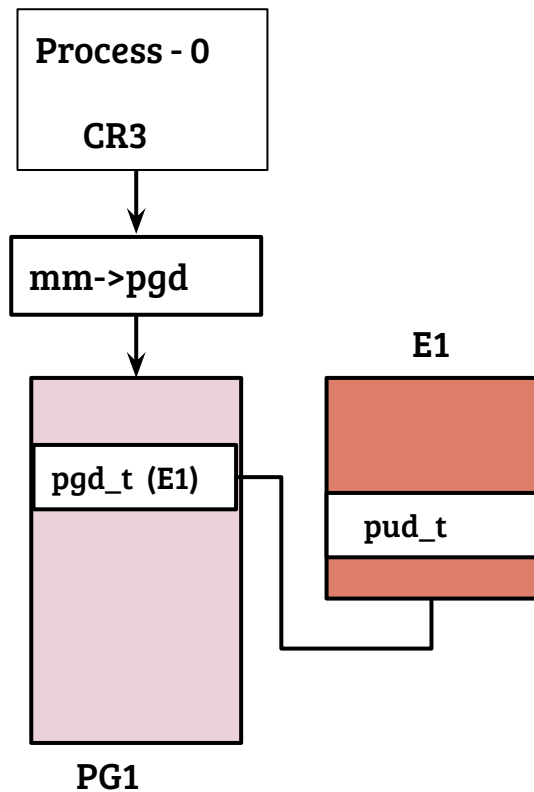
- A child process page table inherits the **kernel mappings** of the parent
- By implication, the inheritance tree is rooted at the first process
- Mapping changes → update mapping in every process?
  - **Does not look good!**

# Linux strives on family values!

- A child process page table inherits the **kernel mappings** of the parent
- By implication, the inheritance tree is rooted at the first process
- Mapping changes → update mapping in every process?
  - Does not look good!


Solution: Every process owns its own **pgd** entries but inherits the kernel **pgd** entries from the parent :-)

# Solution overview




- One (or more) entries in PGD-level (level-4) reserved for kernel mapping
- How many?
- Depends on VA-range covered by one entry and the kernel VA size

# Solution overview




- All updates to E1 are visible across all the processes
- So we are at peace! Not really.

# Virtual memory layout (x86\_64)



- User virtual addresses use the LSB 47 bits
- Kernel virtual address does not start from 0x8000000000, but from 0xffff800000000000
- Why? Because X86 hardware enforces if 47th bit is one, 48-63 must be set to one


# Process address space (user + kernel)



- Virtual address space is split into two parts, user VA and kernel VA
- Kernel mappings are isolated from user through **S/U** bit of page table entry
- Advantages: isolation + efficiency
- What is the need for direct map?



# Process address space (user + kernel)



- Virtual address space is split into two parts, user VA and kernel VA
- Kernel mappings are isolated from user through **S/U** bit of page table entry
- Advantages: isolation + efficiency
- What is the need for direct map? **Helps in mapping physical address to an already mapped kernel vaddr**

# Issue with shared address space

```
char array[256 * 4096];    //__aligned(4k);  
char secret = *(char *) 0xffff888000000000;  
array[secret << 12] = 0;
```

- This program will result in an exception → Segmentation fault
- Everything seems to be under control. What is the problem then?

# Information leakage through out-of-order execution

```
1.  mov RCX, $0xFFFF888000000000;  
2.  mov RBX, $array;  
3.  mov AL, [RCX];
```

```
4.  shl RAX, $0xC;
```

```
5.  mov RBX, qword [RBX + RAX];
```

Executed  
out-of-order

## Exception handler

```
1.  cmp CR2, $userend;  
2.  Jg raise_segv;  
3.  .....  
4.  .....  
5.  raise_segv:  
6.  .....
```

- By the time the instruction in line#3 is committed (and a fault is raised), instructions in line#4 and #5 are completed out-of-order

# Side-effect: access footprint

1. `char array[256 * 4096];` `//__aligned(4k);`
2. `char secret = *(char *) 0xffff888000000000;`
3. `array[secret << 12] = 0;`

Array (before the program execution): block 0 == {0 - 4095} etc.




Array (after out-of-order execution of #3) {assume secret = k}



Accessed

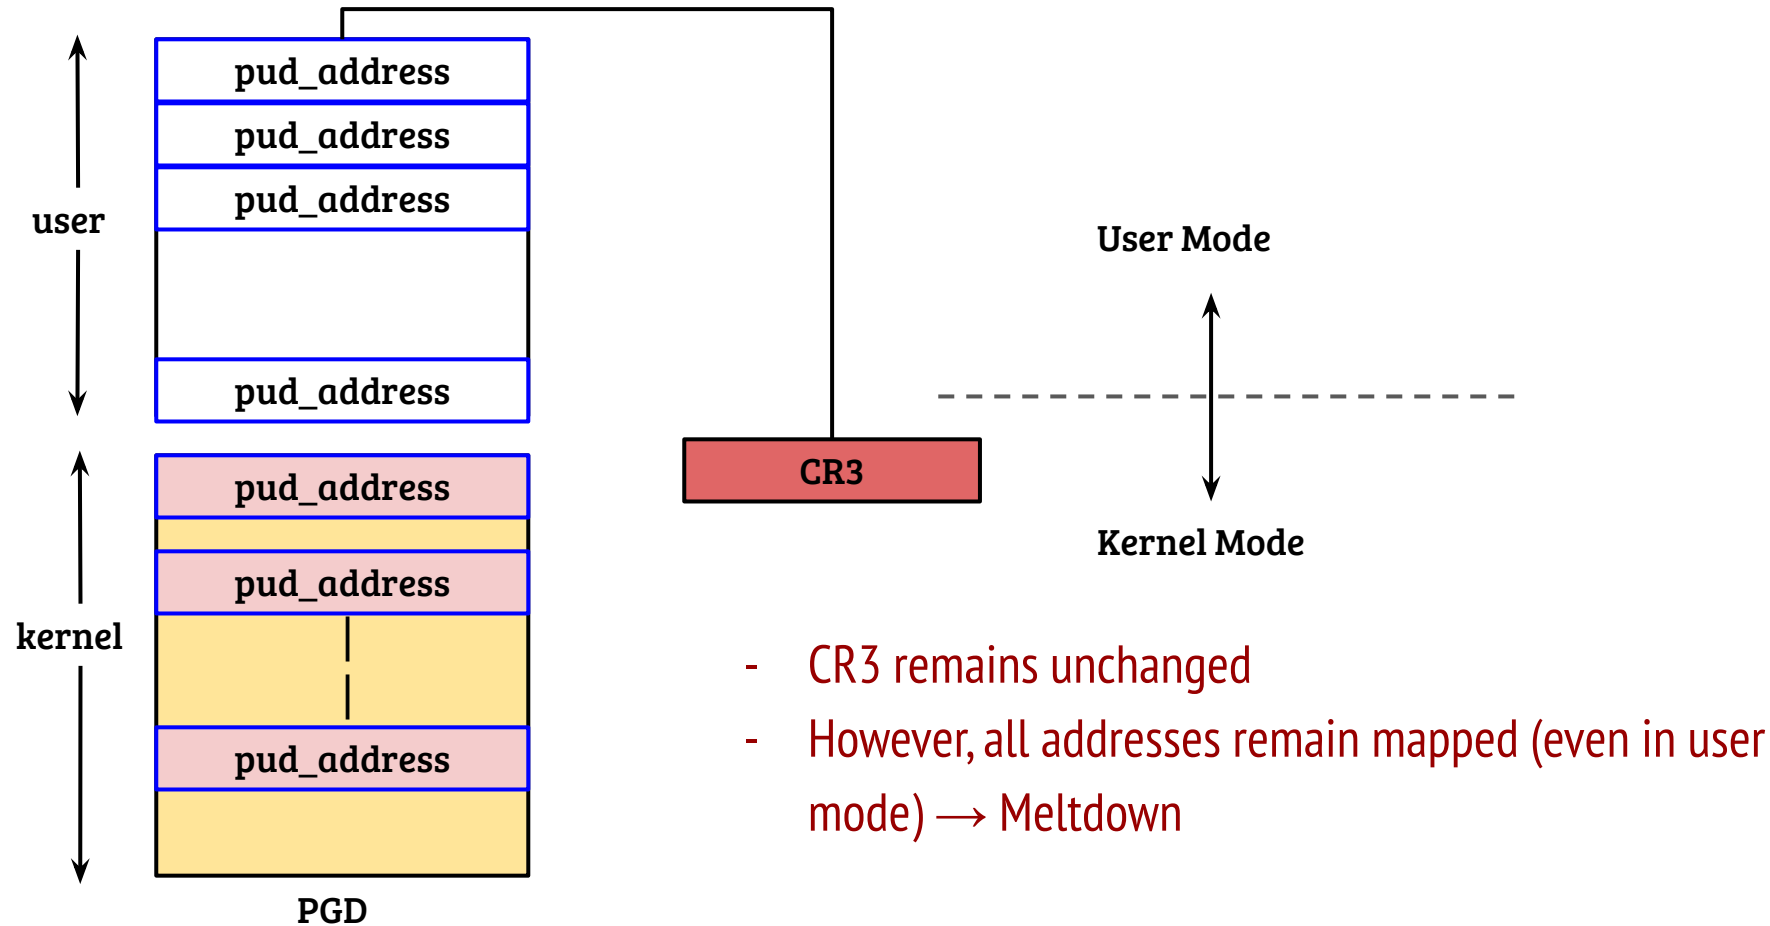
# OOO vulnerability + Flush-Reload

1. unsigned time[256];
2. char array[256 \* 4096];
3. flush\_array(array);
4. char secret = \*(char \*) 0xffff888000000000;
5. array[secret << 12] = 0;
6. for(i=0; i<256; ++i)
7.     access\_and\_time(array, time, i);
8. secret = find\_index\_with\_min\_time( time);




- Result: indirectly read the value of secret
- Meltdown is easy.... Some subtle points still remain
- What is the fix?

# Linux paging (before PTI)



# Linux paging (with PTI)



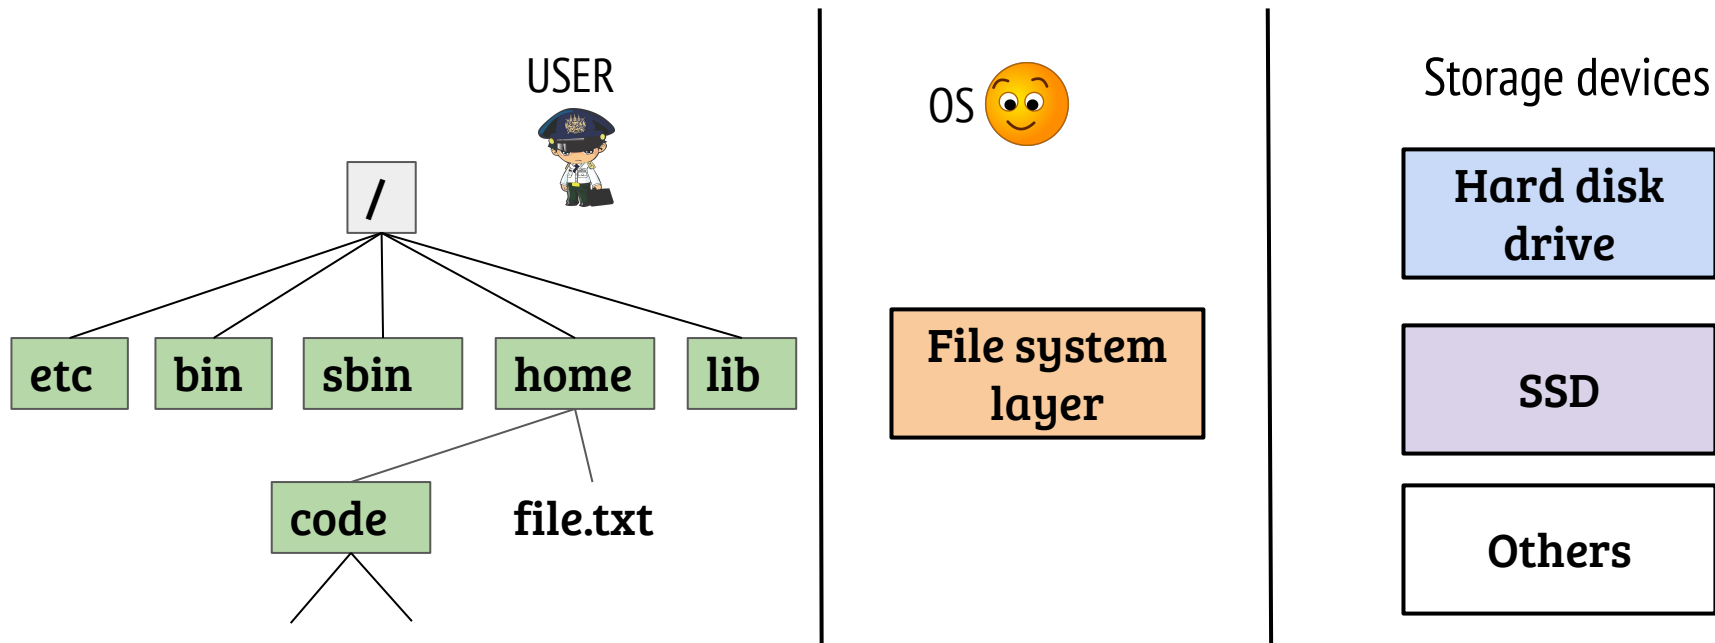
# CS614: Linux Kernel Programming

## File System Overview

Debadatta Mishra, CSE, IIT Kanpur




# Recap: file system




- File system is an important OS subsystem
  - Provides abstractions like files and directories
  - Hides the complexity of underlying storage devices

# File system interfacing




- User process identify files through a file handle a.k.a. file descriptors
- In UNIX, the POSIX file API is used to access files, devices, sockets etc.
- Important file related system calls?

# File system interfacing




- User process identify files through a file handle a.k.a. file descriptors
- In UNIX, the POSIX file API is used to access files, devices, sockets etc.
- Important file related system calls: `open`, `close`, `read`, `write`, `lseek`, `dup`, `stat`, `select`, `poll` ...

# Process view of file



- Per-process file descriptor table with pointer to a “file” object
- file object → inode is many-to-one

# Process view of file



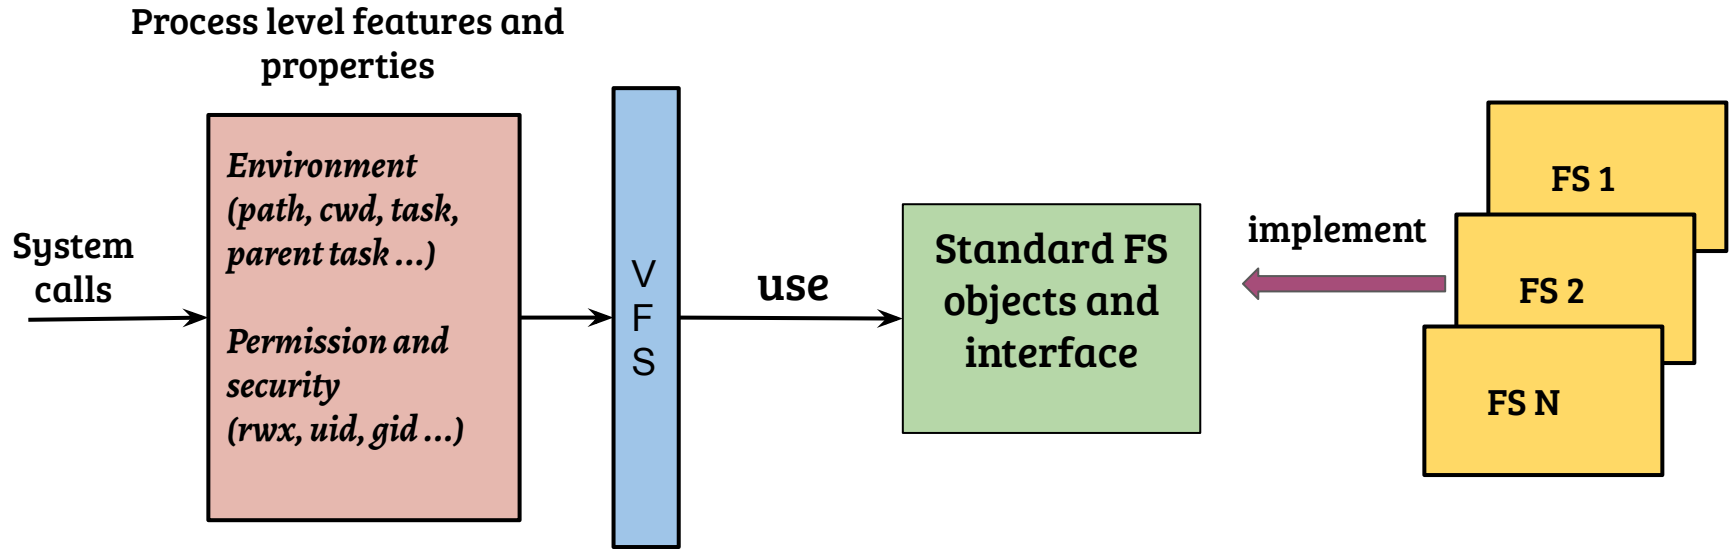
- What happens to the FD table and the file objects across `fork()`?
  - What happens in `exec()`?
- Can multiple FDs point to the same file object?

# Process view of file

PCB (P1)

- What happens to the FD table and the file objects across `fork( )`?
  - What happens in `exec( )`?
- The FD table is copied across `fork( )`  $\Rightarrow$  File objects are shared
- On `exec`, open files remain shared by default
- Can multiple FDs point to the same file object?
- Yes, duped FDs share the same file object (within a process)

# Linux virtual file system (VFS)



- Object and interface choices guided by API requirement (mostly)
- Sometimes standards (e.g., POSIX) determines the interfacing
- Implementation can be different for different file systems


# Linux virtual file system (VFS)

Process level features and  
properties

- VFS to Disk, How the dots are connected?
  - How a FS is created?
  - How system calls are mapped to the file system?
- Object and interface choices guided by API requirement (mostly)
- Sometimes standards (e.g., POSIX) determines the interfacing
- Implementation can be different for different file systems




# Step-1: Disk device partitioning



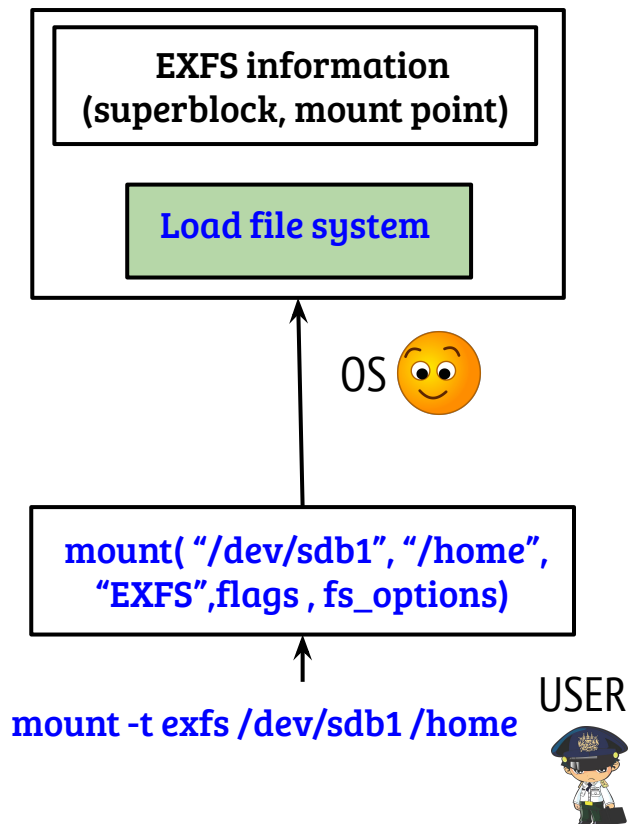
- Partition information is stored in the boot sector of the disk
- Creation of partition is the first step
  - It does not create a file system
- A file system is created on a partition to manage the physical device and present the logical view
- All file systems provide utilities to initialize file system on the partition (e.g., MKFS)

# Step 2: File system creation



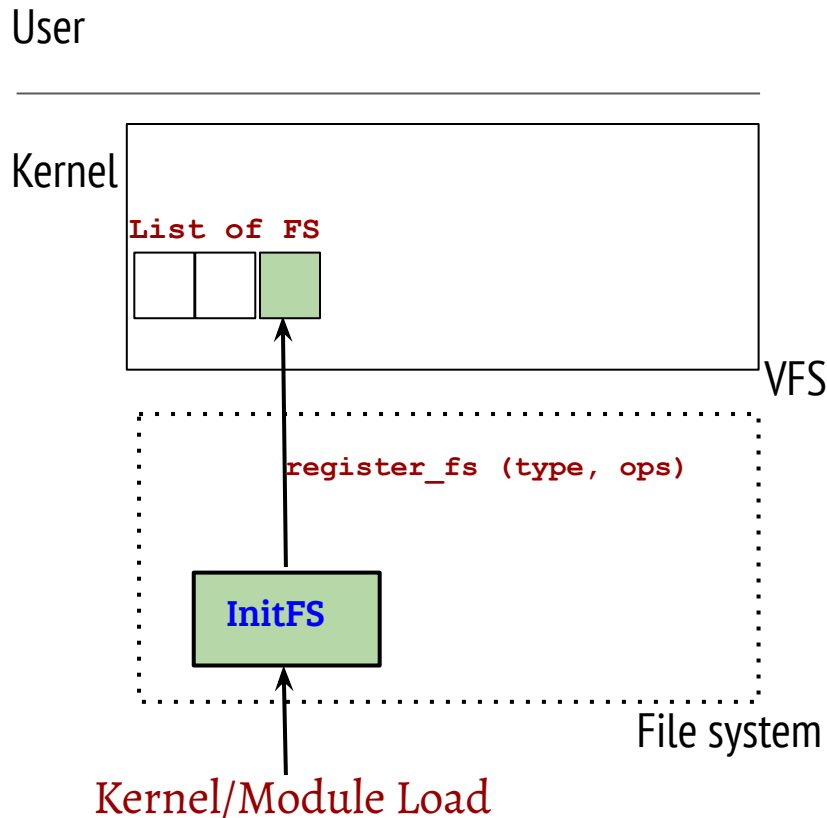
- MKFS creates initial structures in the logical partition
  - Creates the entry point to the filesystem (known as the super block)
  - At this point the file system is ready to be mounted

# Step 3: File system mounting



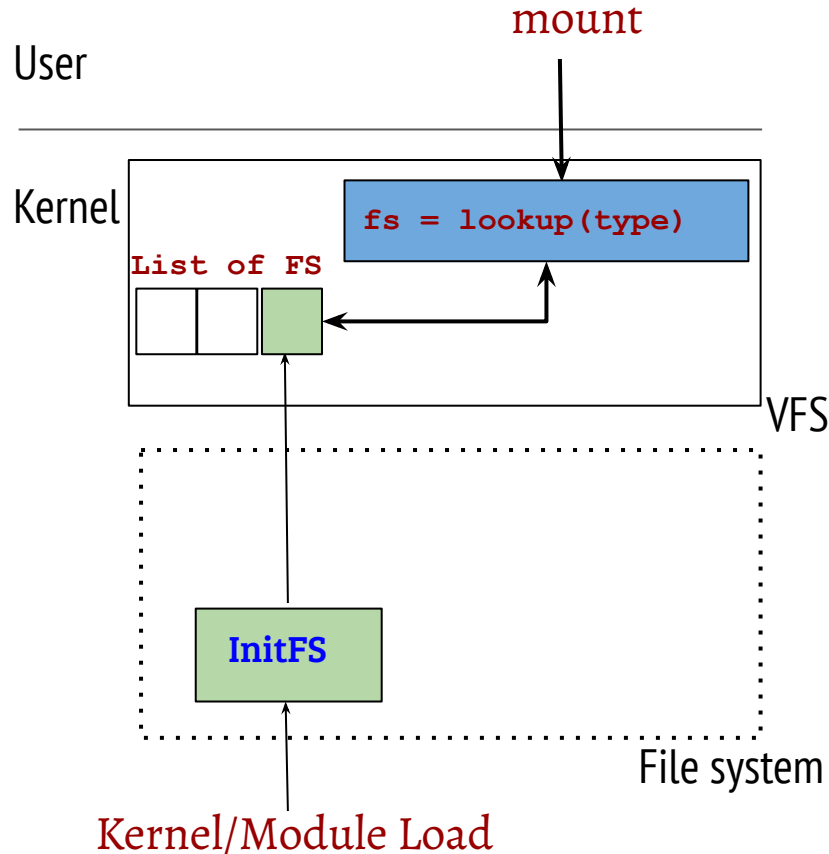
- *mount( )* associates a superblock with the file system mount point
- Example: The OS will use the superblock associated with the mount point “/home” to reach any file/dir under “/home”
- Superblock is a copy of the on-disk superblock along with other information

# Details of FS mount in Linux (simplified)



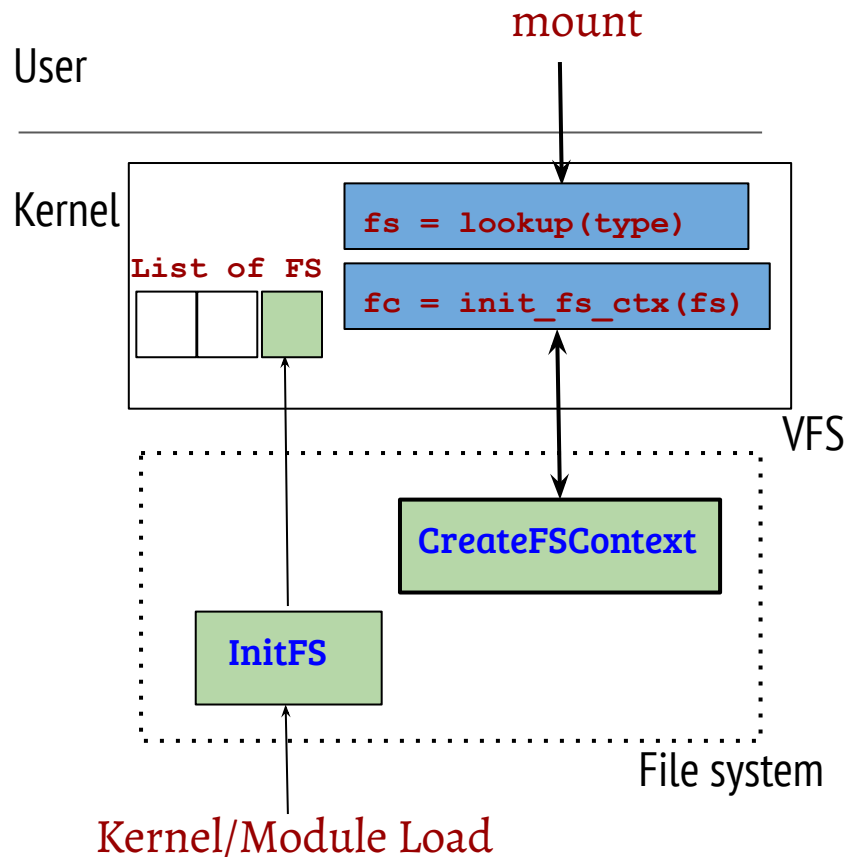
- File system registers itself with the VFS layer during initialization
  - “type” is the identity of the file system (e.g., ext4)
  - “ops” contains the callbacks for different events such as context initialization and mount
- VFS layer maintains a list of registered file system types

# Details of FS mount in Linux (simplified)



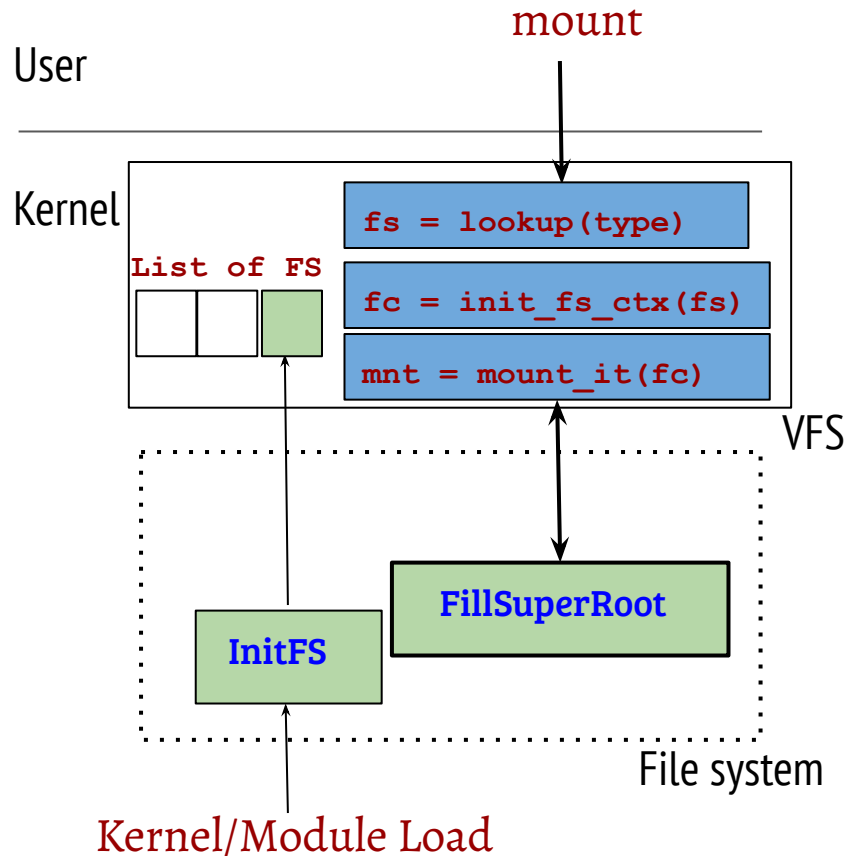
- System call handler for mount looks up the FS type

# Details of FS mount in Linux (simplified)



- System call handler for mount looks up the FS type
- Creates a context – an instance of the FS for a given mount point

# Details of FS mount in Linux (simplified)



- System call handler for mount looks up the FS type
- Creates a context – an instance of the FS for a given mount point
- The FS fills superblock and root inode information (by performing disk block I/O)
- A new mount point is created at the VFS layer for future use. What kind of use?

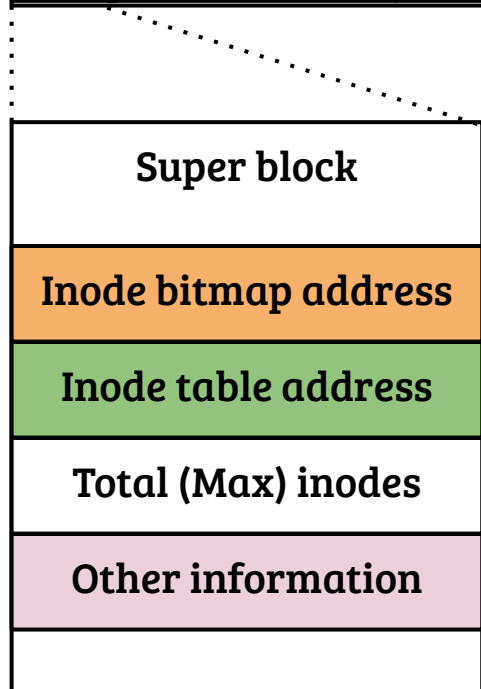

# Structure of an example superblock

```
struct superblock{  
    u16 block_size;  
    u64 num_blocks;  
    u64 last_mount_time;  
    u64 root_inode_num;  
    u64 max_inodes;  
    disk_off_t inode_table;  
    disk_off_t blk_usage_bitmap;  
    ...  
};
```

- Superblock contains information regarding the device and the file system organization in the disk
- Pointers to different metadata related to the file system are also maintained by the superblock
  - Ex: List of free blocks is required before adding data to a new file/directory



# Typical file system organization (on-disk)



- Given any inode number, load the inode structure into memory

```
inode_t *get_inode(SB *sb, long ino){  
    inode_t *inode = alloc_mem_inode();  
    read_disk(inode, sb -> inode_table +  
                ino * sizeof(inode));  
    return inode;  
}
```

# File system organization



- File system is mounted, the inode number for root of the file system (i.e., the mount point) is known, root inode can be accessed. However,
- How to search/lookup files/directories under root inode?
- Specifically,
  - How to locate the content in disk?
  - How to keep track of size, permissions etc.?

*return inode;*


}

# Inode

- A on-disk structure containing information regarding files/directories in the unix systems
  - Represented by a unique number in the file system (e.g., in Linux, “ls -li filename” can be used to print the inode)
  - Contains access permissions, access time, file size etc.
  - *Most importantly, inode contains information regarding the file data location on the device*
- Directory inodes also contain information regarding its content, albeit the content is structured (for searching files)

# Ext2 file system indexing

Ext2/3 inode



Direct pointers {PTR [0] to PTR [11]}



File block address (0 -11)

Single indirect {PTR [12]}



File block address (12 -1035)

Double indirect {PTR [13]}



File block address (1036 to 1049611)

Triple indirect {PTR [14]}



File block address (?? to ??)

# File system organization

SB	Block bitmaps	Inode bitmaps	Inode table	Data blocks
----	---------------	---------------	-------------	-------------

- File system is mounted, the inode number for root of the file system (mount point) is known, root inode can be accessed. However,
- How to search/lookup files/directories under root inode?
- Specifically,
  - How to locate the content in disk?
  - Index structures in inode are used to map file offset to disk location
  - How to keep track of size, permissions etc.?
  - Inode is used to maintain these information

# Organizing the directory content

## Fixed size directory entry

```
struct dir_entry{  
    inode_t inode_num;  
    char name[FNAME_MAX];  
};
```

- Fixed size directory entry is a simple way to organize directory content
- Advantages: avoid fragmentation
- Disadvantages: space wastage

# Flat organization of directory entries

## Fixed size directory entry

```
struct dir_entry{  
    inode_t inode_num;  
    char name[FNAME_MAX];  
};
```

## Variable size directory entry

```
struct dir_entry{  
    inode_t inode_num;  
    u8 entry_len;  
    char name[name_len];  
};
```

- Variable sized directory entries contain length explicitly
- Advantages: less space wastage (compact)
- Disadvantages: fragmentation issues

# File system organization

- File system is mounted, the inode number for root of the file system (mount point) is known, root inode can be accessed. However,
- How to search/lookup files/directories under root inode?
- Read the content of the root inode and search the next level dir/file
- Specifically,
  - How to locate the content in disk?
  - Index structures in inode are used to map file offset to disk location
  - How to keep track of size, permissions etc.?
  - Inode is used to maintain these information



# File system and caching

- Accessing data and metadata from disk impacts performance
- Many file operations require multiple block access

# File system and caching

- Accessing data and metadata from disk impacts performance
- Many file operations require multiple block access
- Examples:
  - Opening a file

```
fd = open("/home/user/test.c", O_RDWR);
```

# File system and caching

- Accessing data and metadata from disk impacts performance
- Many file operations require multiple block access
- Examples:
  - Opening a file

```
fd = open("/home/user/test.c", O_RDWR);
```

- Normal shell operations

```
/home/user$ ls
```

# File system and caching

- Accessing data and metadata from disk impacts performance
- Many file operations require multiple block access
- Executables, configuration files, library etc. are accessed frequently
- Many directories containing executables, configuration files are also accessed very frequently. Metadata blocks storing inodes, indirect block pointers are also accessed frequently

*/home/user\$ ls*


# File system and caching

- Accessing data and metadata from disk impacts performance
- Can we store frequently accessed disk data in memory?
  - What is the storage and lookup mechanism? Are the data and metadata caching mechanisms same?
  - Are there any complications because of caching?
  - How the cache managed? What should be the eviction policy?

*/home/user\$ ls*

# Block layer caching


## Cached I/O



- Lookup memory cache using the block number as the key
- How does the scheme work for data and metadata?

# Block layer caching


## Cached I/O



- Lookup memory cache using the block number as the key
- How does the scheme work for data and metadata?
- For data caching, file offset to block address mapping is required before using the cache

# Block layer caching

## Cached I/O




- Lookup memory cache using the block number as the key
- How does the scheme work for data and metadata?
- For data caching, file offset to block address mapping is required before using the cache
- Works fine for metadata as they are addressed using block numbers




# File layer caching (Linux page cache)

## Cached I/O



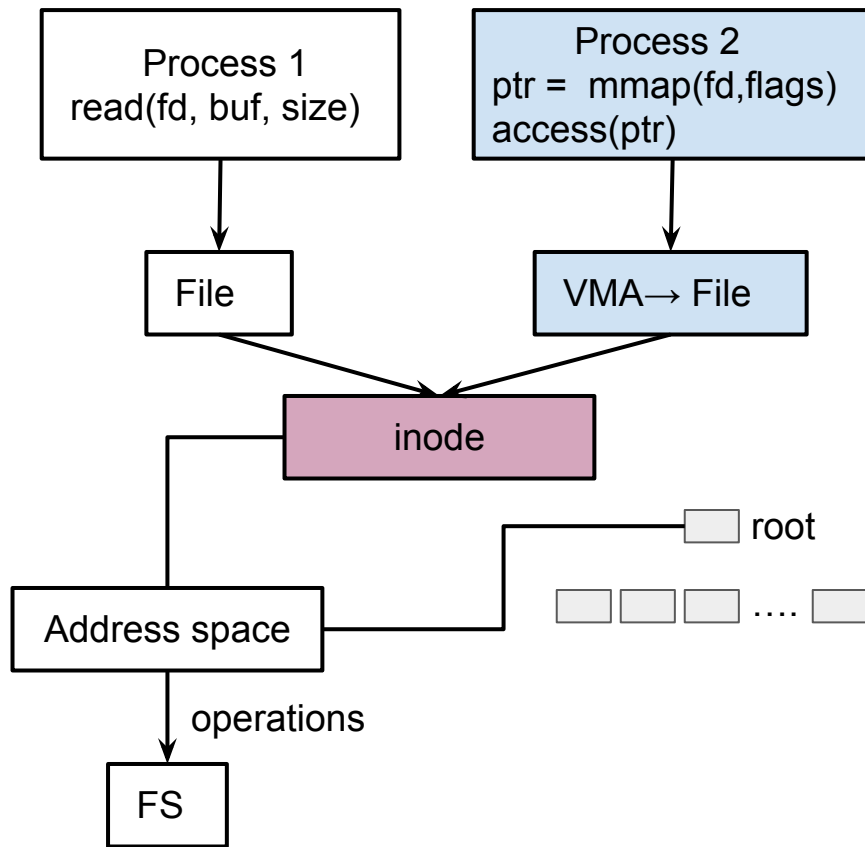
- Store and lookup memory cache using {inode number, file offset} as the key
- For data, index translation is not required for file access
- Metadata may not have a file association, should be handled differently (using a special inode may be!)

# Linux page cache: A multi-purpose FS caching layer



- Requirement: File block lookup at different offsets
  - File size can range from very small to huge
- Recall: mmap-ing a file creates a VMA struct
- Should handle both file I/O and page faults

# File → Inode → Address spaces → Page Cache



- A per inode cache
  - Lookup, insert, evict, dirty-flush
- Radix tree
  - Root pointed by address space struct
  - Operations at a page size (4K) granularity
- Homework: For a given file, find the number of file blocks cached in PC

# File system and caching

- Accessing data and metadata from disk impacts performance
- Can we store frequently accessed disk data in memory?
  - What is the storage and lookup mechanism? Are the data and metadata caching mechanisms same?
  - File layer caching is desirable as it avoids index accesses on hit, special mechanism required for metadata.
  - Are there any complications because of caching?
  - How the cache managed? What should be the eviction policy?

# Caching and consistency

- Caching may result in inconsistency, but what type of consistency?

# Caching and consistency

- Caching may result in inconsistency, but what type of consistency?
- System call level guarantees
  - Example-1: If a write( ) system call is successful, data must be written
  - Example-2: If a file creation is successful then, file is created.
  - Difficult to achieve with asynchronous I/O

# Caching and consistency

- Caching may result in inconsistency, but what type of consistency?
- System call level guarantees
  - Example-1: If a write( ) system call is successful, data must be written
  - Example-2: If a file creation is successful then, file is created.
  - Difficult to achieve with asynchronous I/O
- Consistency w.r.t. file system invariants
  - Example-1: If a block is pointed to by an inode data pointers then, corresponding block bitmap must be set
  - Example-2: Directory entry contains an inode, inode must be valid
  - Possible, require special techniques

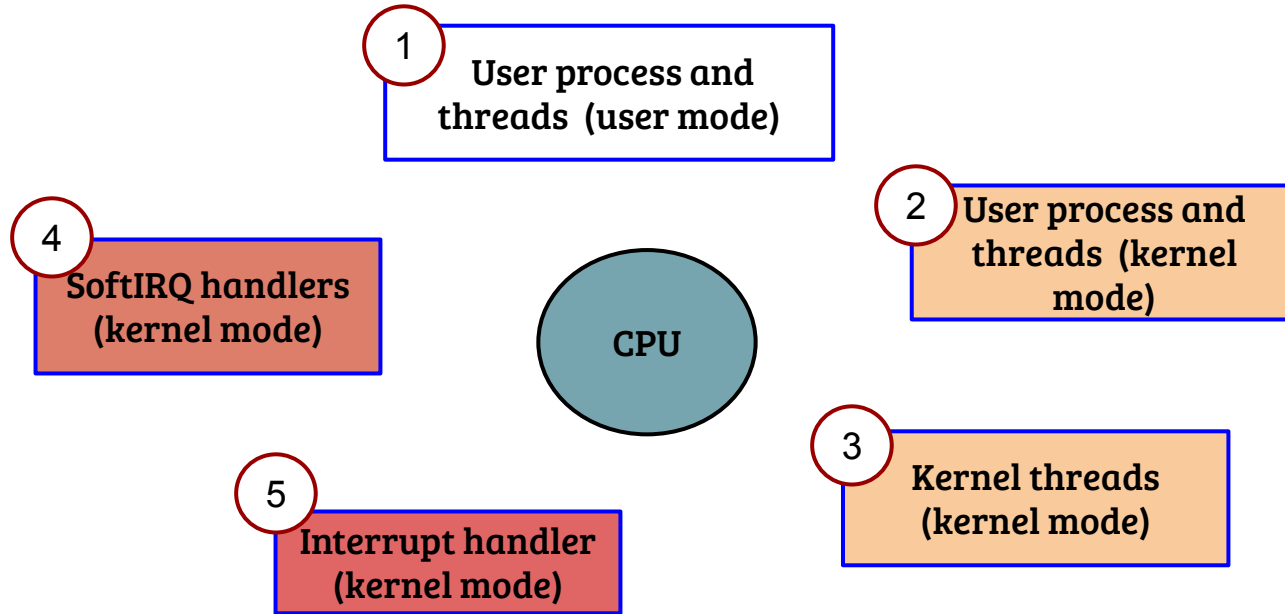
# CS614: Linux Kernel Programming

## Process, Thread, Kernel Threads

Debadatta Mishra, CSE, IIT Kanpur




# Recap: Execution contexts in Linux



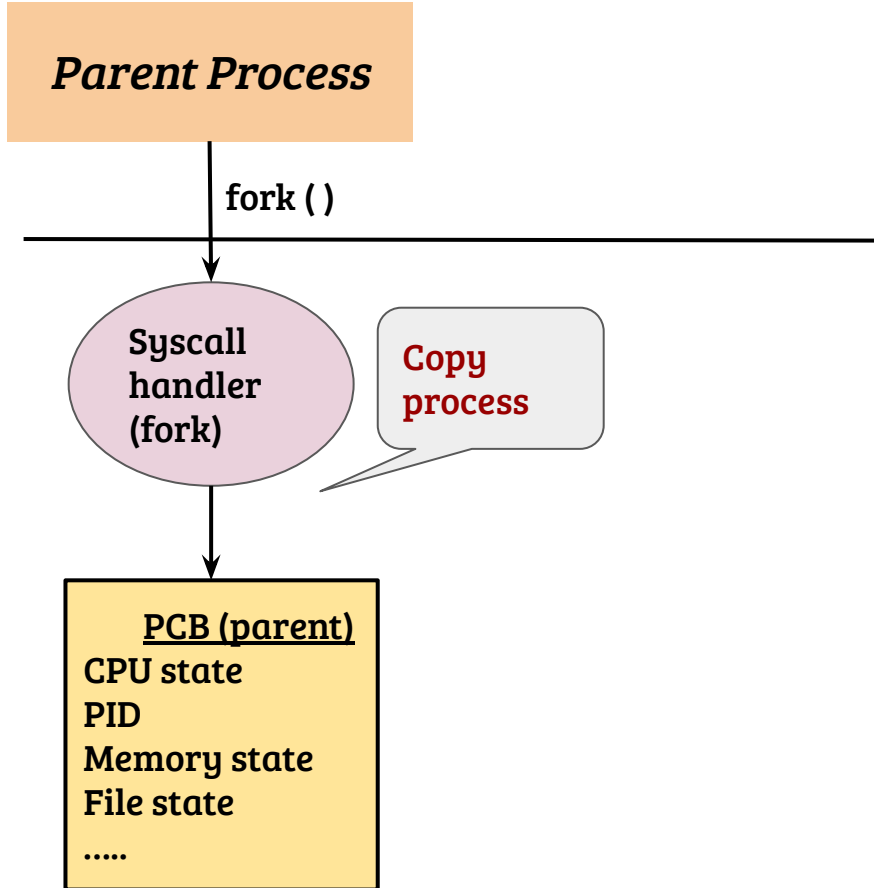
- In a linux system, the CPU can be executing in one of the above contexts
- For (3), (4) and (5), the context is not associated with any user process

# Process creation - fork( )

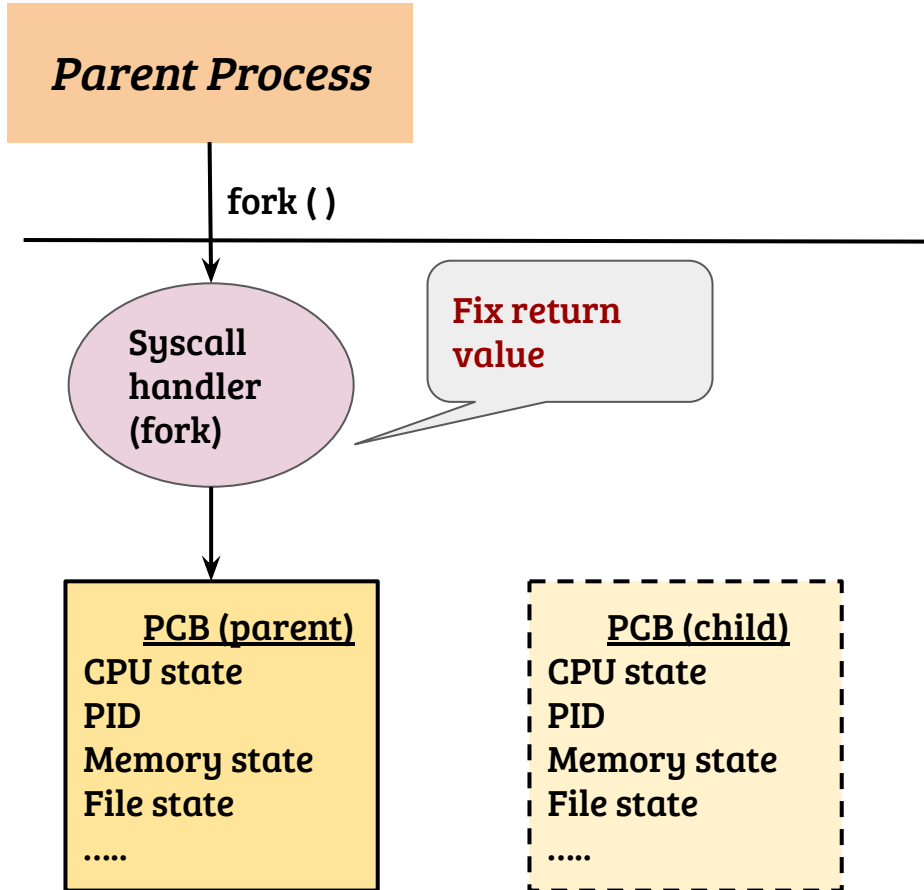


- fork( ) system call is weird; not a typical “privileged” function call
- fork( ) creates a new process; a *duplicate* of calling process
- On success, fork
  - Returns PID of child process to the caller (parent)
  - Returns 0 to the child

# Typical implementation of fork

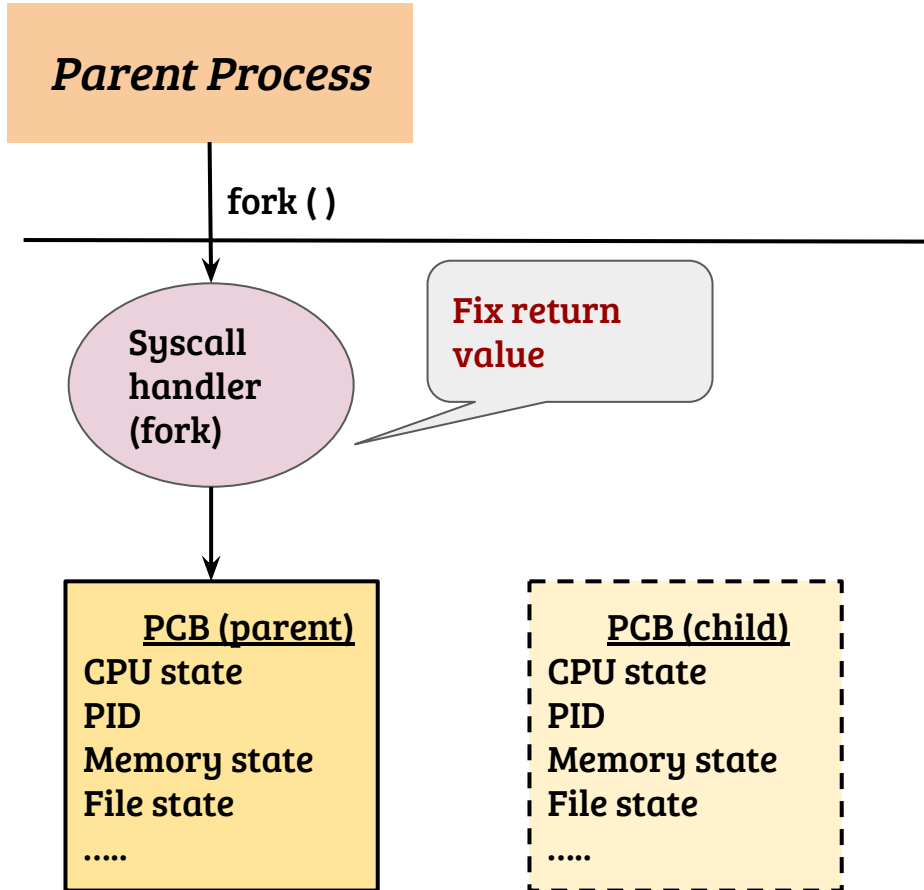


# Typical implementation of fork



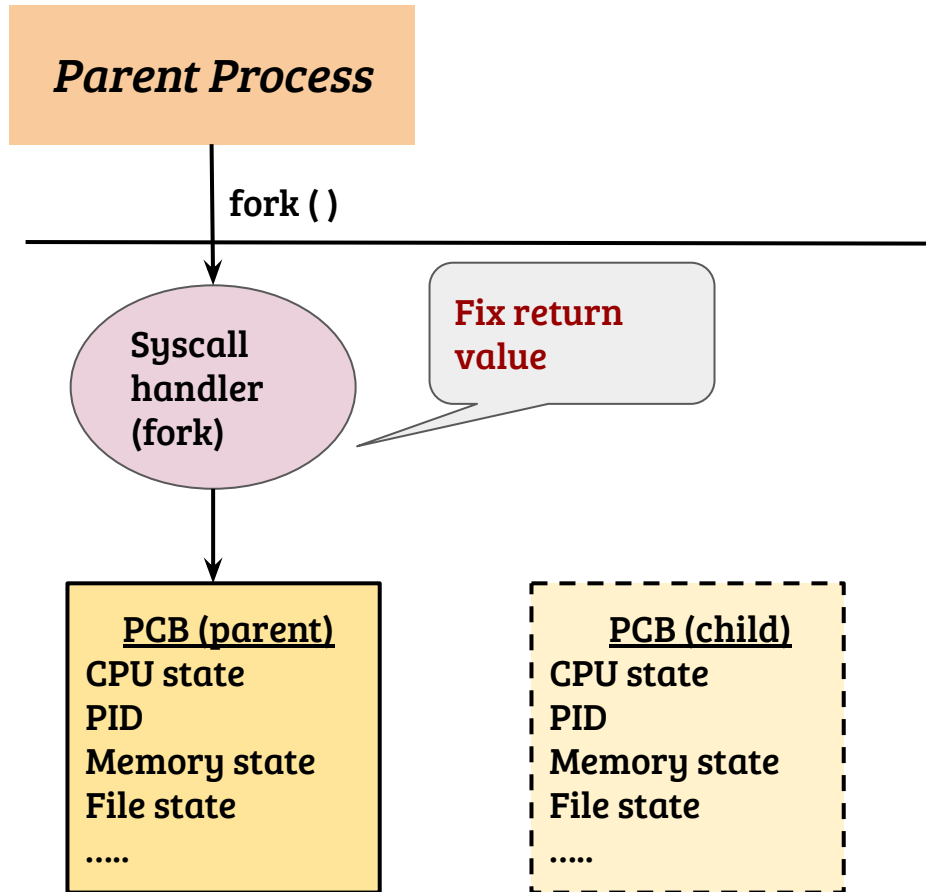
- Child should get '0' and parent gets PID of child as return value. How?

# Typical implementation of fork



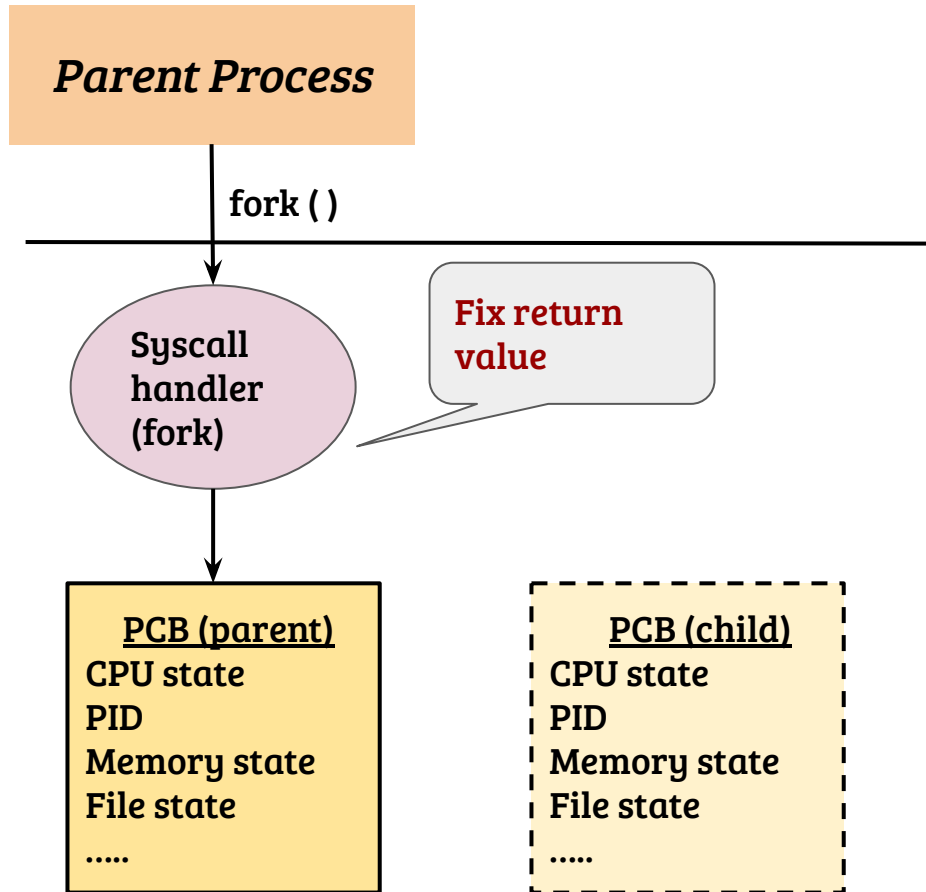
- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child

# Typical implementation of fork



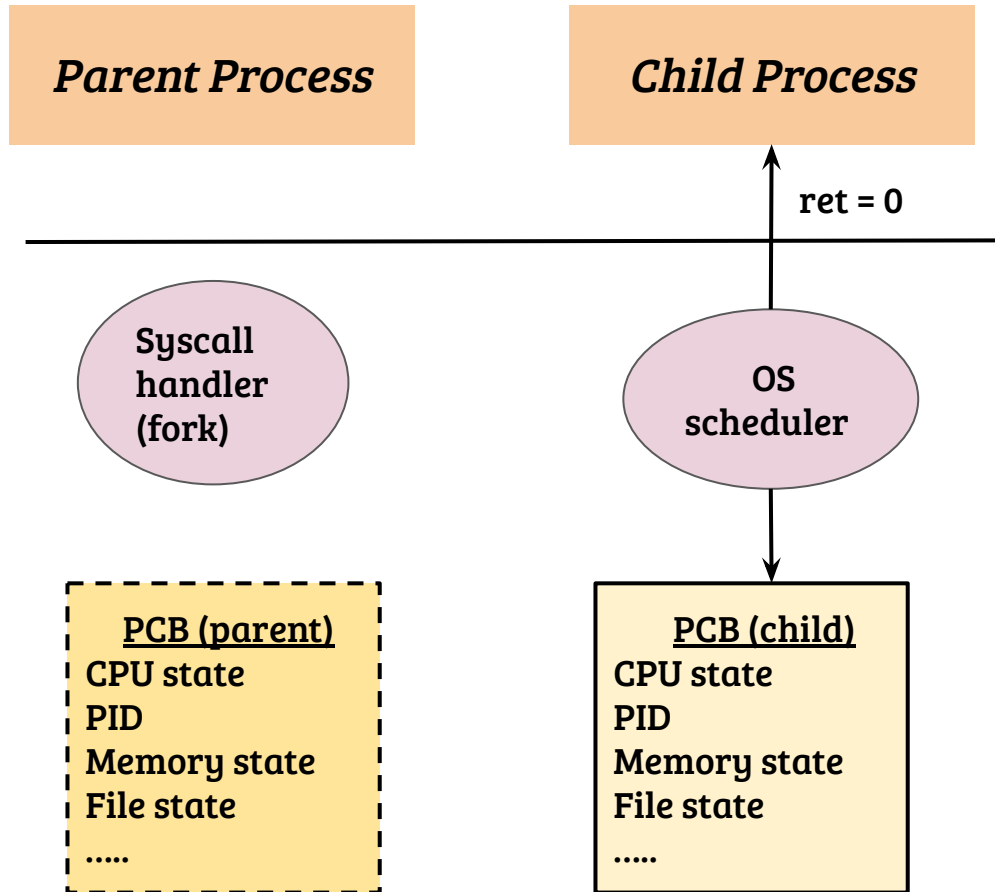
- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?

# Typical implementation of fork



- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?
- When OS schedules the child process

# Typical implementation of fork



- PC is next instruction after `fork( )` syscall, for both parent and child
- Child memory is an exact copy of parent
- Parent and child diverge from this point

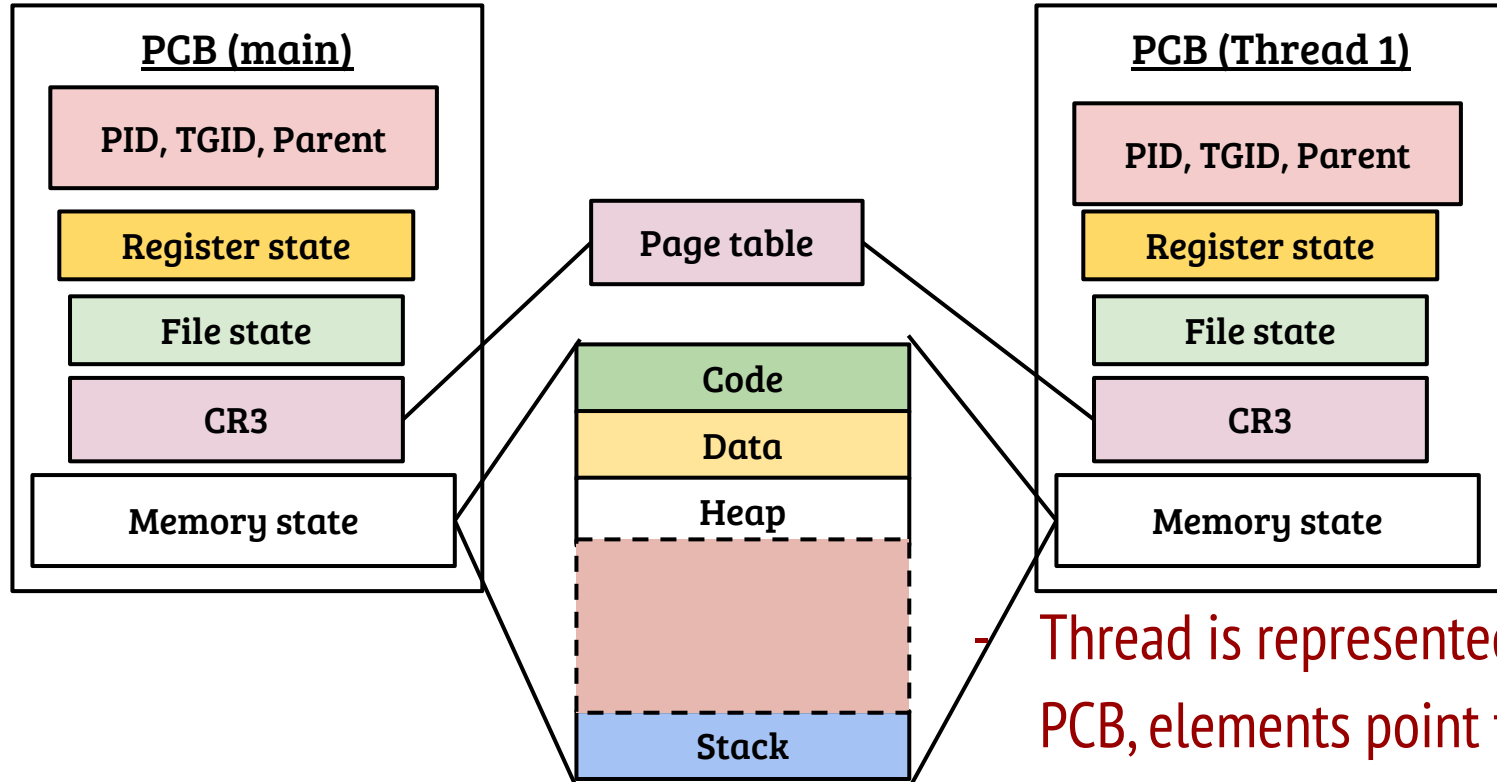


# User threads using posix thread API

```
int pthread_create( pthread_t *tid, pthread_attr_t *attr,  
                  void * (*thfunc) (void*), void *arg);
```

- Creates a thread with “tid” as its handle and the thread starts executing the function pointed to by the “thfunc” argument
- A single argument (of type void \*) can be passed to the thread
- Thread attribute can be used to control the thread behavior e.g., stack size, stack address etc. Passing NULL sets the defaults
- Returns 0 on success.
- Thread termination: return from thfunc, pthread\_exit( ) or pthread\_cancel( )
- In Linux, pthread\_create and fork implemented using clone( ) system call

# PCB of a multithreaded process (Linux)



- Thread is represented by a separate PCB, elements point to the structure containing subsystem level info.

# The clone system call

```
int clone(int (*fn)(void *), void *child_stack, int flags, void *arg, ...)
```

- Parent can control the execution of new process (execution and stack)
- Provides flexibility to the parent to share parts of its execution context in a selective manner
- Examples flags:
  - CLONE\_FILES: Share files between parent and new process
  - CLONE\_VM: Share the address space
  - CLONE\_VFORK: Execution of parent process is suspended

# Clone: Implementation in Linux kernel

- Syscall handler for clone should provide flexible sharing. Implementation?

# Clone: Implementation in Linux kernel

- Syscall handler for clone should provide flexible sharing. Implementation?
  - Syscall Handler → Kernel clone → Copy process
  - Depending on flags, different subsystems are copied or shared
- Depending on the usage, the saved user state is required to be changed.

Why? How implemented?

# Clone: Implementation in Linux kernel

- Syscall handler for clone should provide flexible sharing. Implementation?
  - Syscall Handler → Kernel clone → Copy process
  - Depending on flags, different subsystems are copied or shared
- Depending on the usage, the saved user state is required to be changed.

Why? How implemented?

- For pthreads, the SP and RIP need to be changed
- Change the register states during CPU thread copy
- Changes to the kernel space of newly created execution context required.

Why? How implemented?

# Clone: Implementation in Linux kernel

- Syscall handler for clone should provide flexible sharing. Implementation?
  - Syscall Handler → Kernel clone → Copy process
  - Depending on flags, different subsystems are copied or shared
- Depending on the usage, the saved user state is required to be changed.


Why? How implemented?

- For pthreads, the SP and RIP need to be changed
- Change the register states during CPU thread copy
- Changes to the kernel space of newly created execution context required.

Why? How implemented?

- Child can not return in the same path, returns through a special stub


# Load a new binary - `exec( )`



- Replace the calling process by a new executable
  - Code, data etc. are replaced by the new process
  - Usually, open files remain open

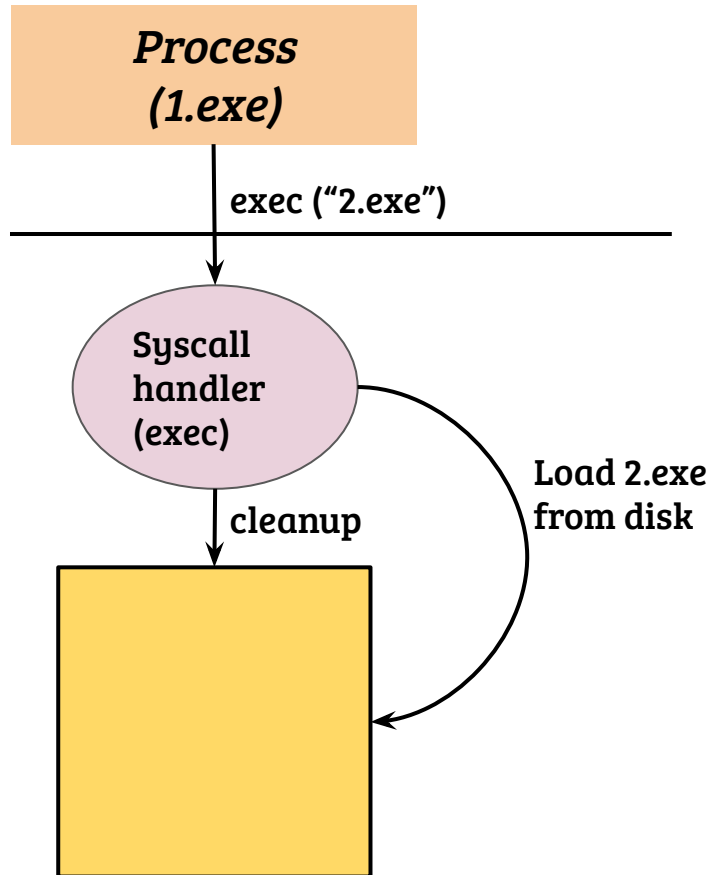


# Typical implementation of exec




- The calling process commits self destruction! (almost)

# Typical implementation of exec



- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same

# Typical implementation of exec



- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same
- On return, new executable starts execution
- PC is loaded with the starting address of the newly loaded binary

# Exec: Implementation in Linux kernel

- When should the self destruction of address space take place? What are the design choices?

# Exec: Implementation in Linux kernel

- When should the self destruction of address space take place? What are the design choices?
  - Can not destroy until validity is checked; validity check not complete until the binary/arguments are examined
  - Duplicated processing vs. working with a fresh (discardable) space
  - There would be a point of no return, delayed is better
- How does the kernel parse the binary (and deduce entry address)? What about command line arguments?

# Exec: Implementation in Linux kernel

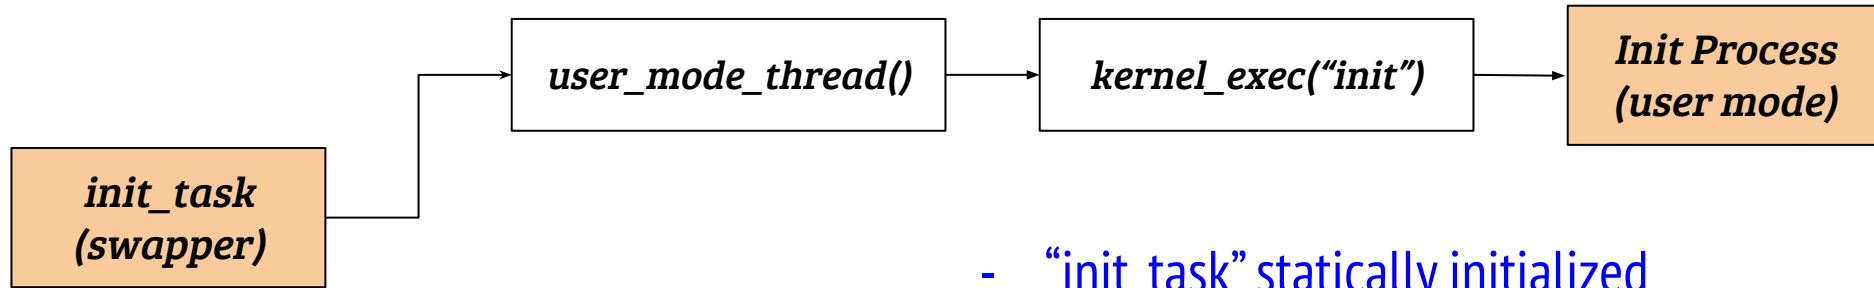
- When should the self destruction of address space take place? What are the design choices?
  - Can not destroy until validity is checked; validity check not complete until the binary/arguments are examined
  - Duplicated processing vs. working with a fresh (discardable) space
  - There would be a point of no return, delayed is better
- How does the kernel parse the binary (and deduce entry address)? What about command line arguments?
  - Basic binary parsing for ELF (and other types) e.g., `load_elf_binary ( )`
  - Command line arguments are placed in the stack

# The first process

- What is the first execution entity in Linux?

# The first process

- What is the first execution entity in Linux?




- “init\_task” statically initialized
- A special “clone” call from the kernel mode to create a thread of execution in kernel till actual init is executed
- Executes user space init based on configuration and default paths



# The first process

- What is the first execution entity in Linux?



- “kthreadd” acts as a kernel thread manager and parent of all kernel threads
- Thread creation list
  - Add a request on kthread create (all types of kernel threads)
  - Wakeup kthreadd
  - Kthreadd → kernel\_thread( )